

2025 AGI Mizzou Hackathon

Team t[AI]ger roar

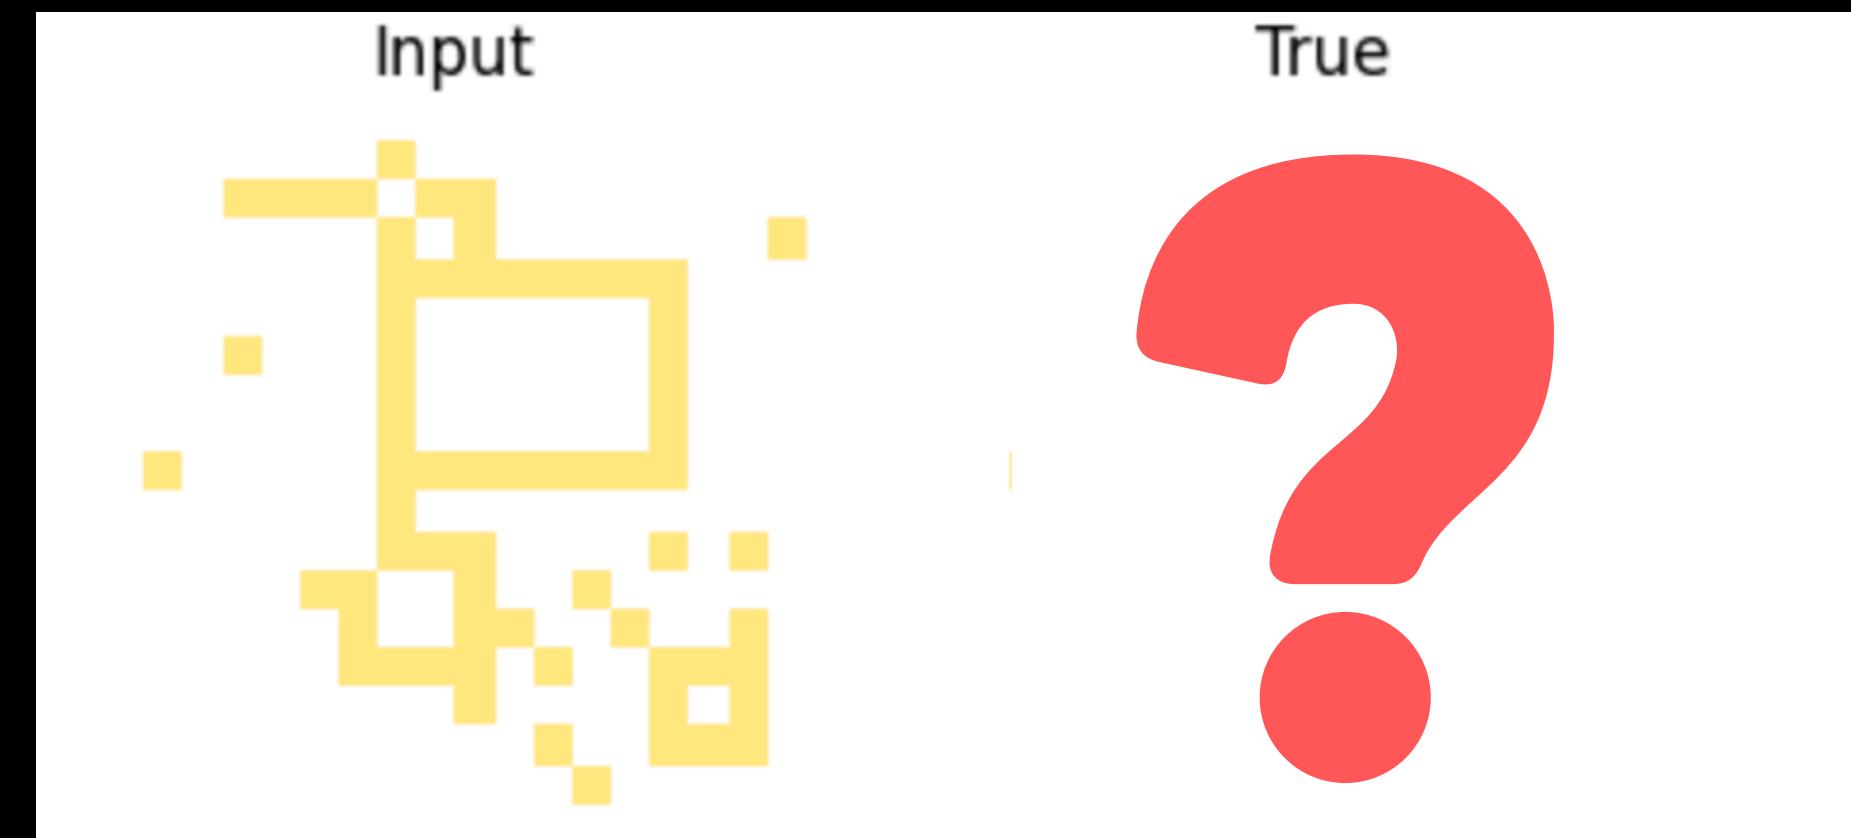
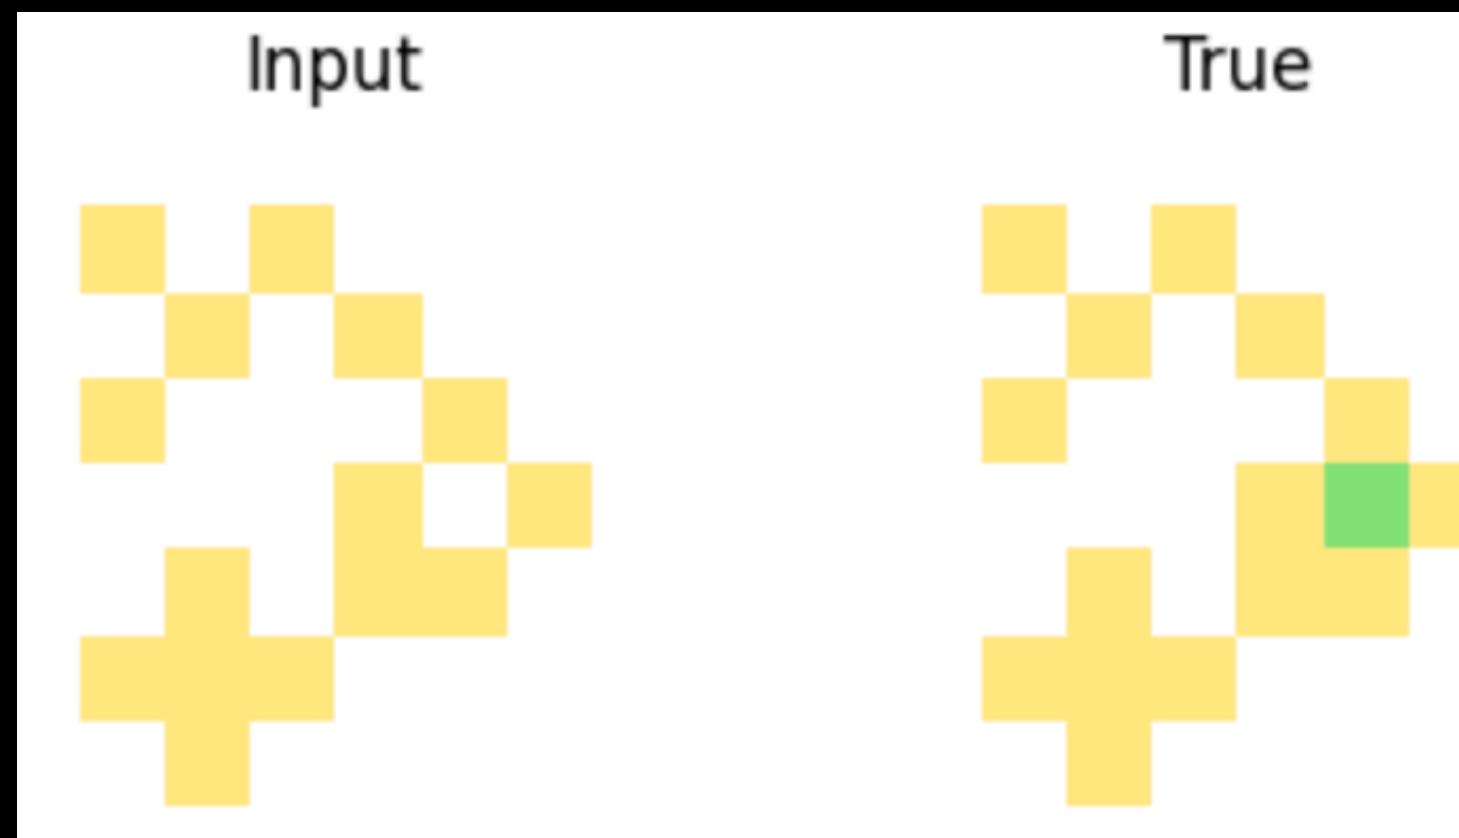
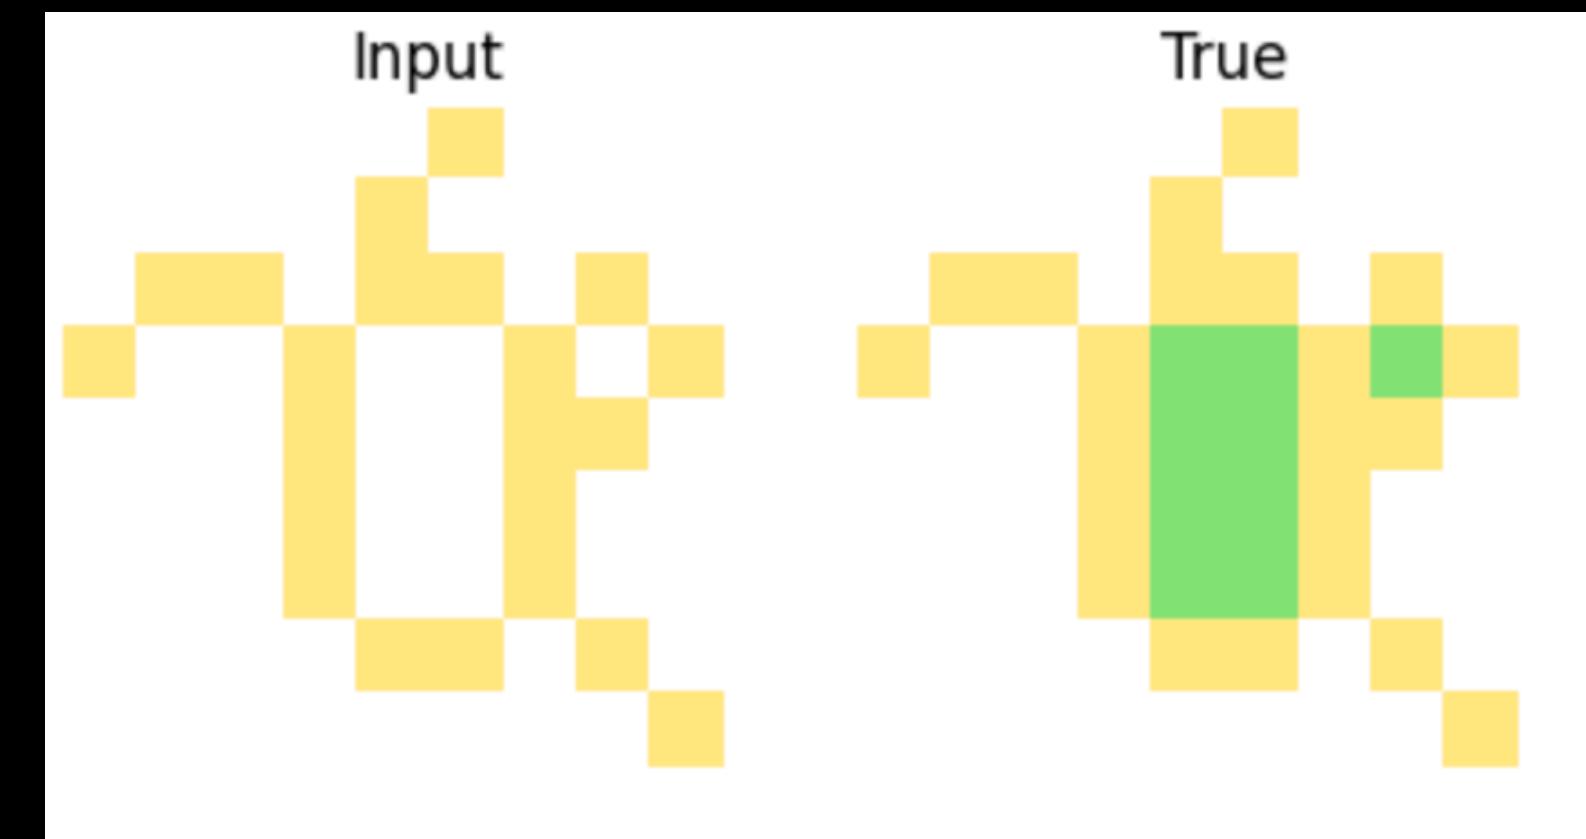
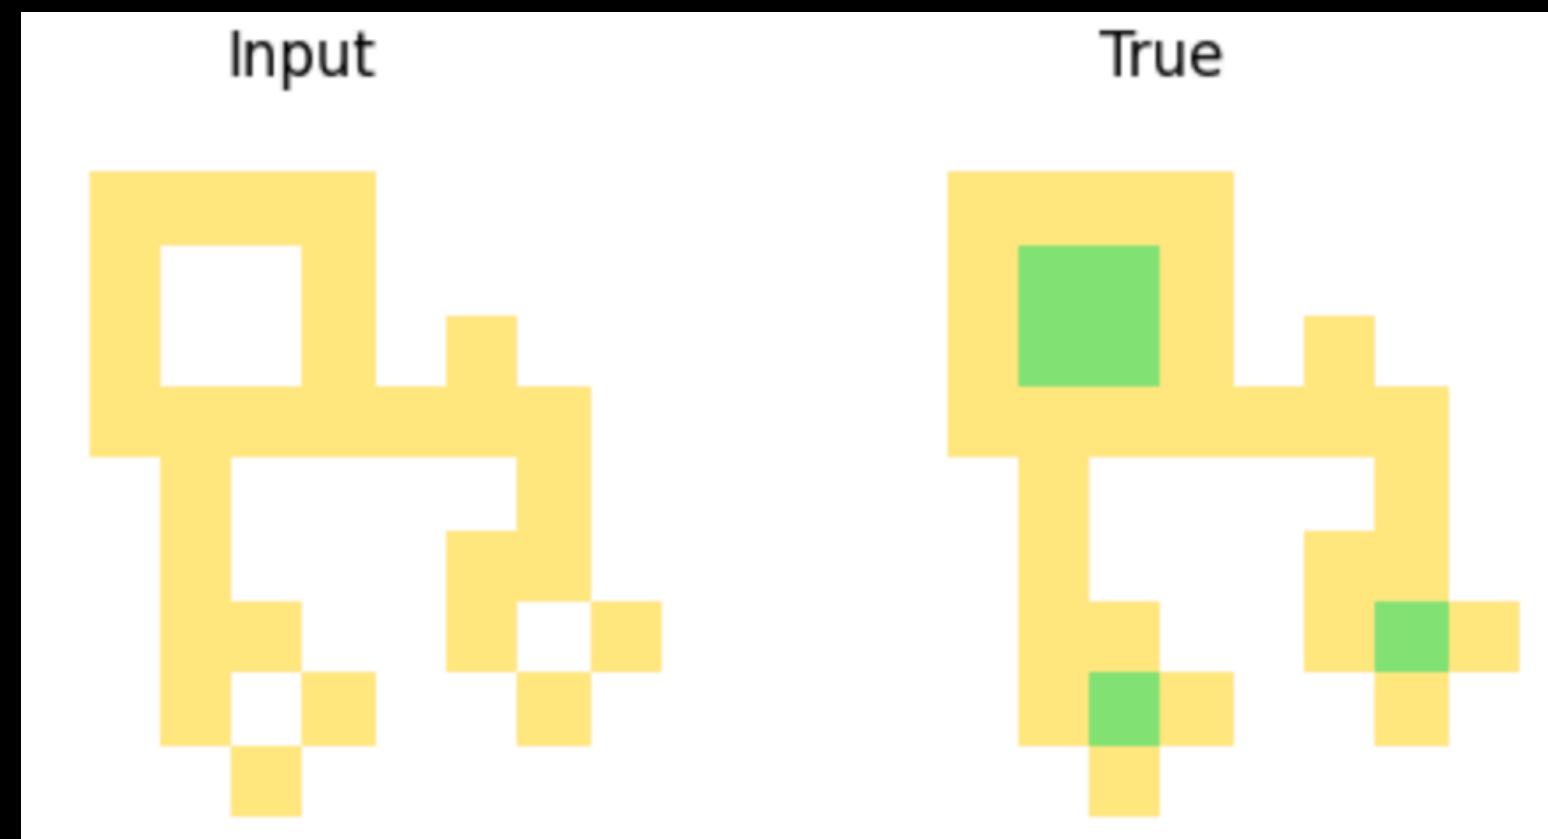
Wen-Hsin Chen | Jeong Wook Lee | Alina Rohulia | Samrat Kumar Dey | Kun-Yi Chen



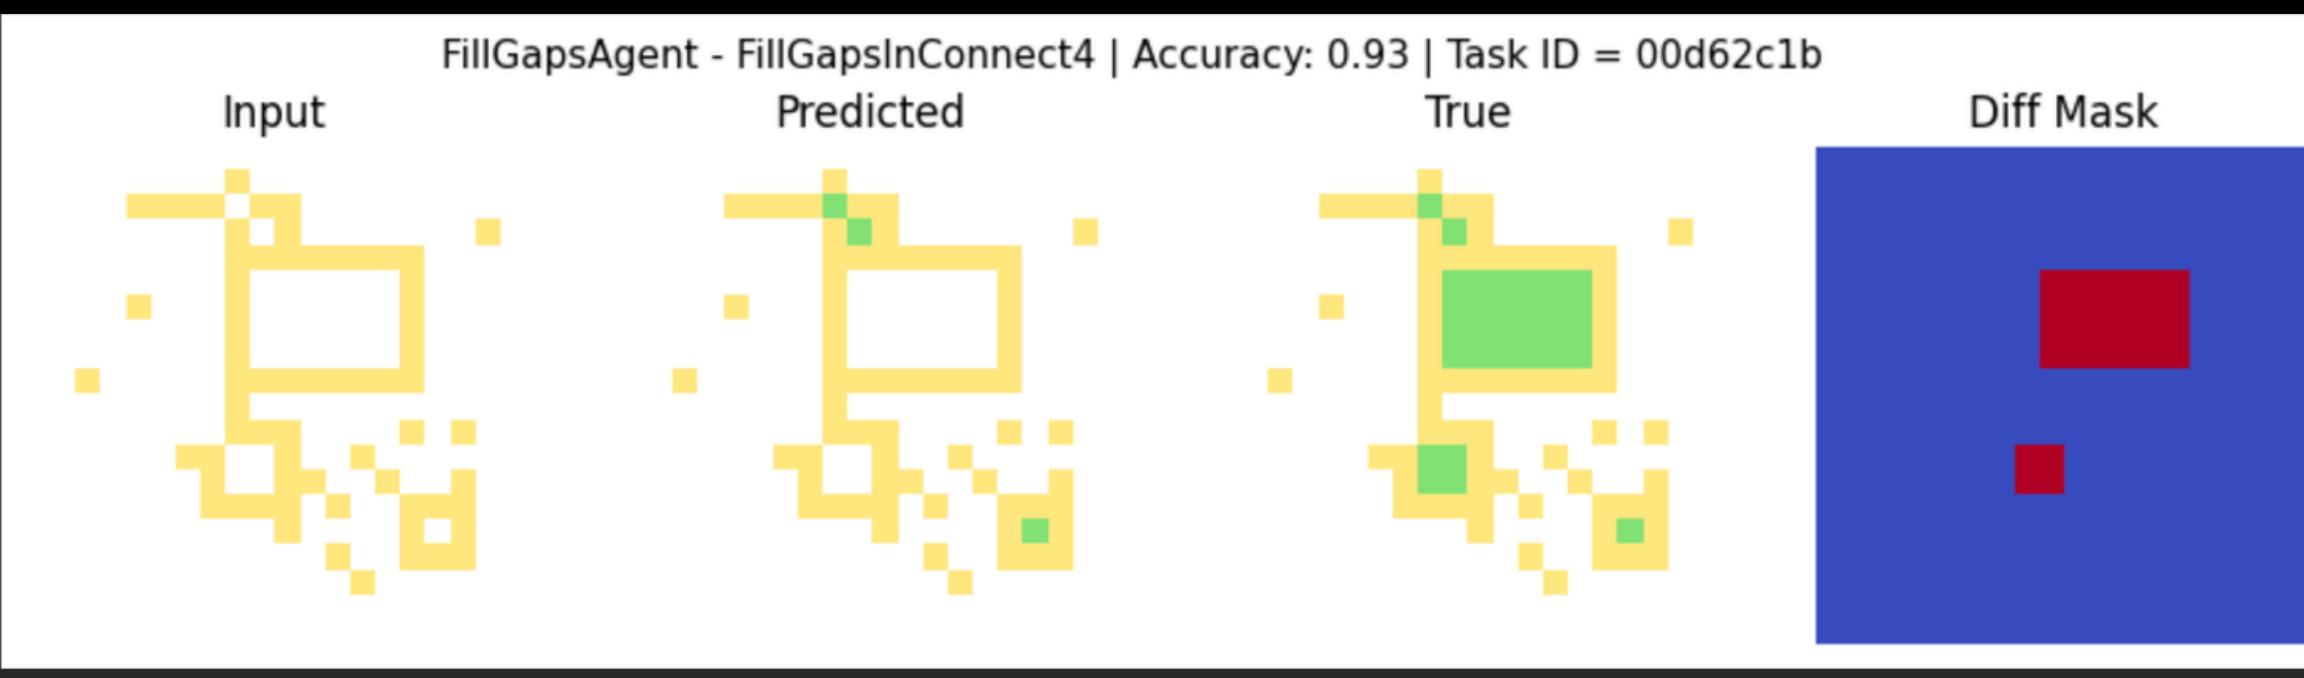
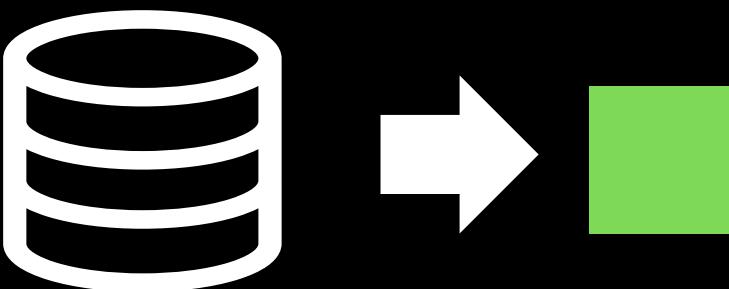
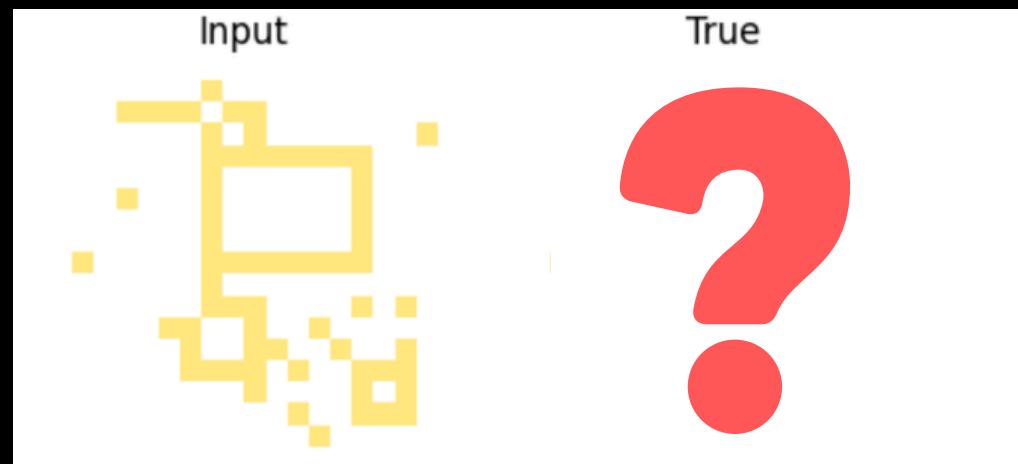
What is ARC-AGI?

**“It’s easy for humans,
but hard for AI.”**

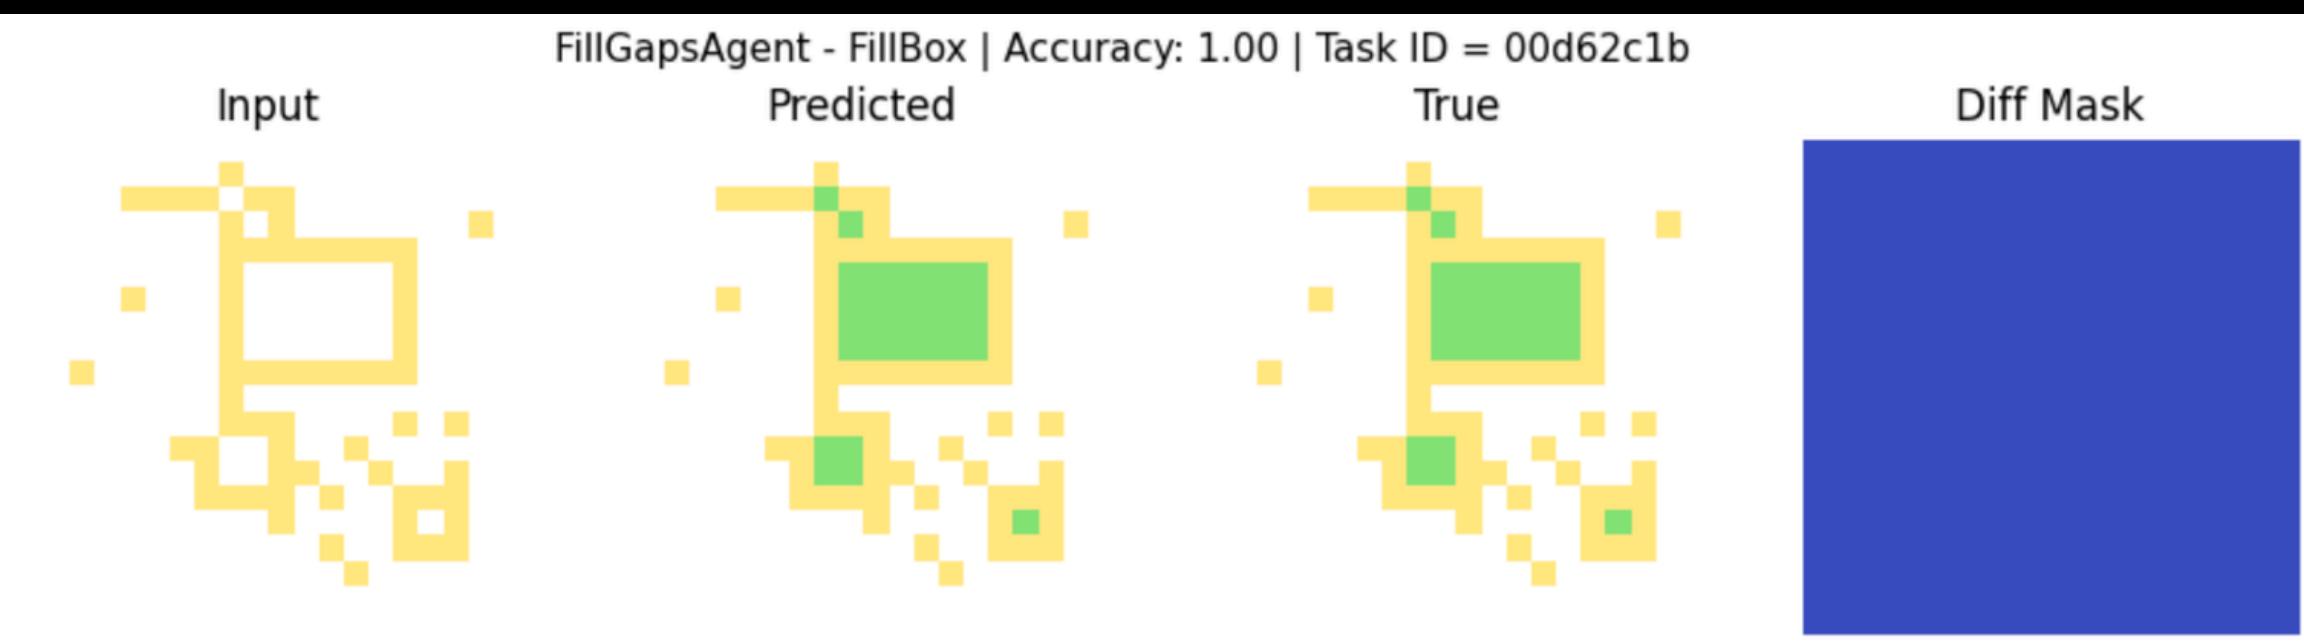
Human solving...



Machine solving...

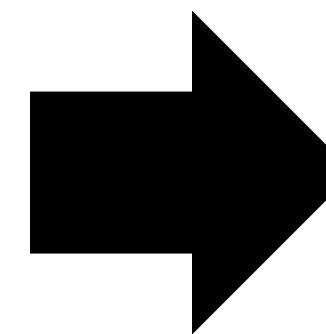
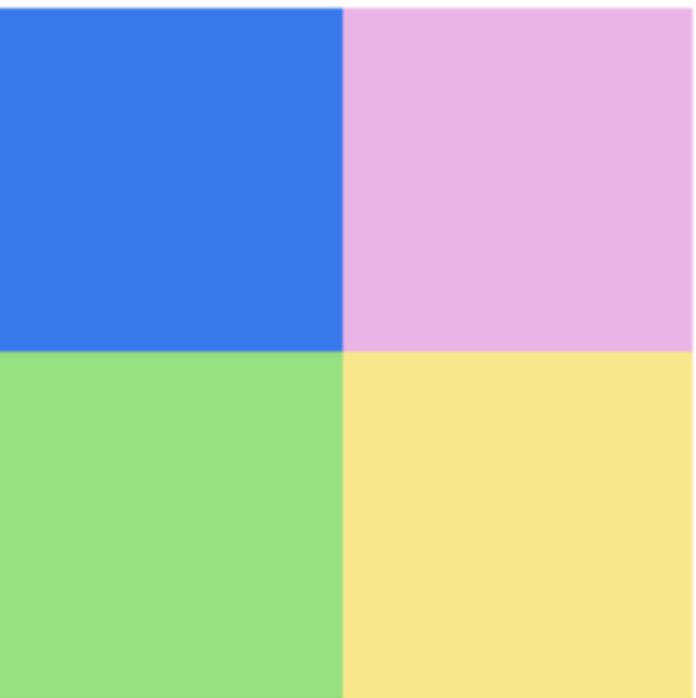


Fill in when $\leftarrow, \uparrow, \rightarrow, \downarrow$ are all **yellow**

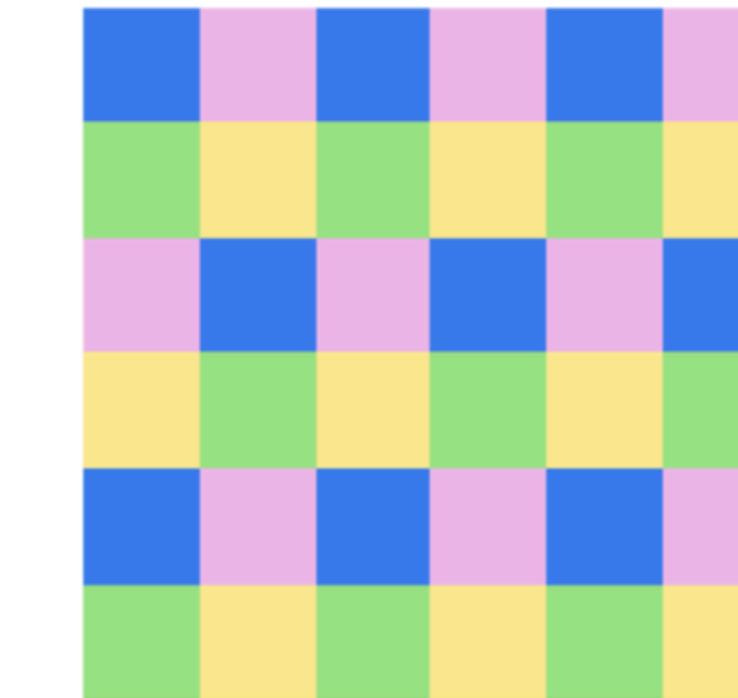


Fill in when box borders are all **yellow**

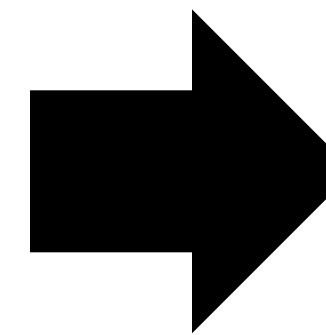
Input



True



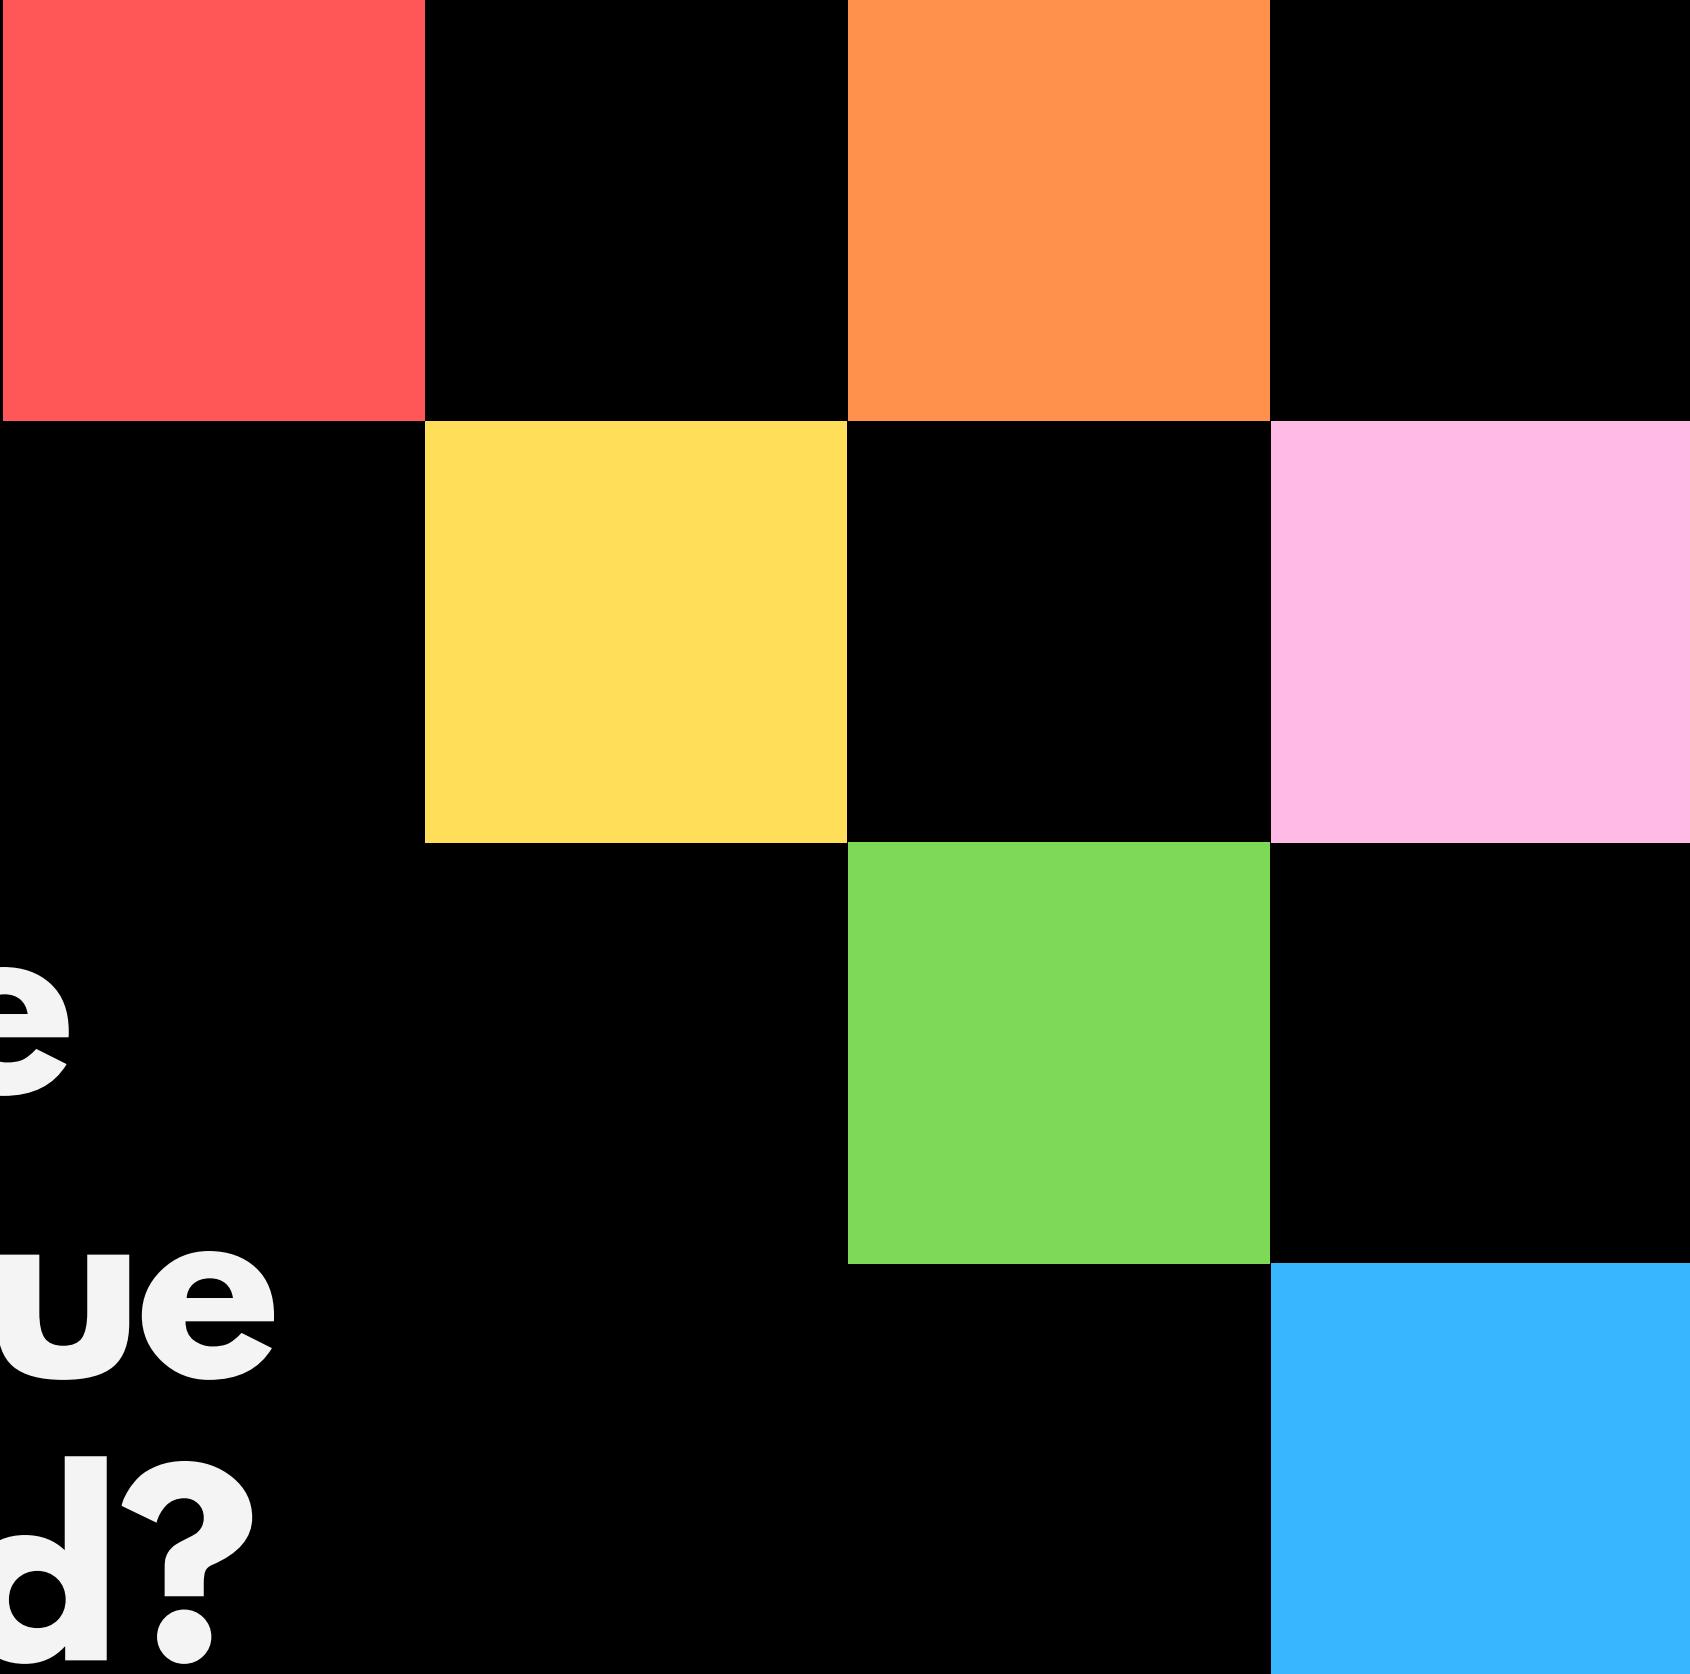
Input



True

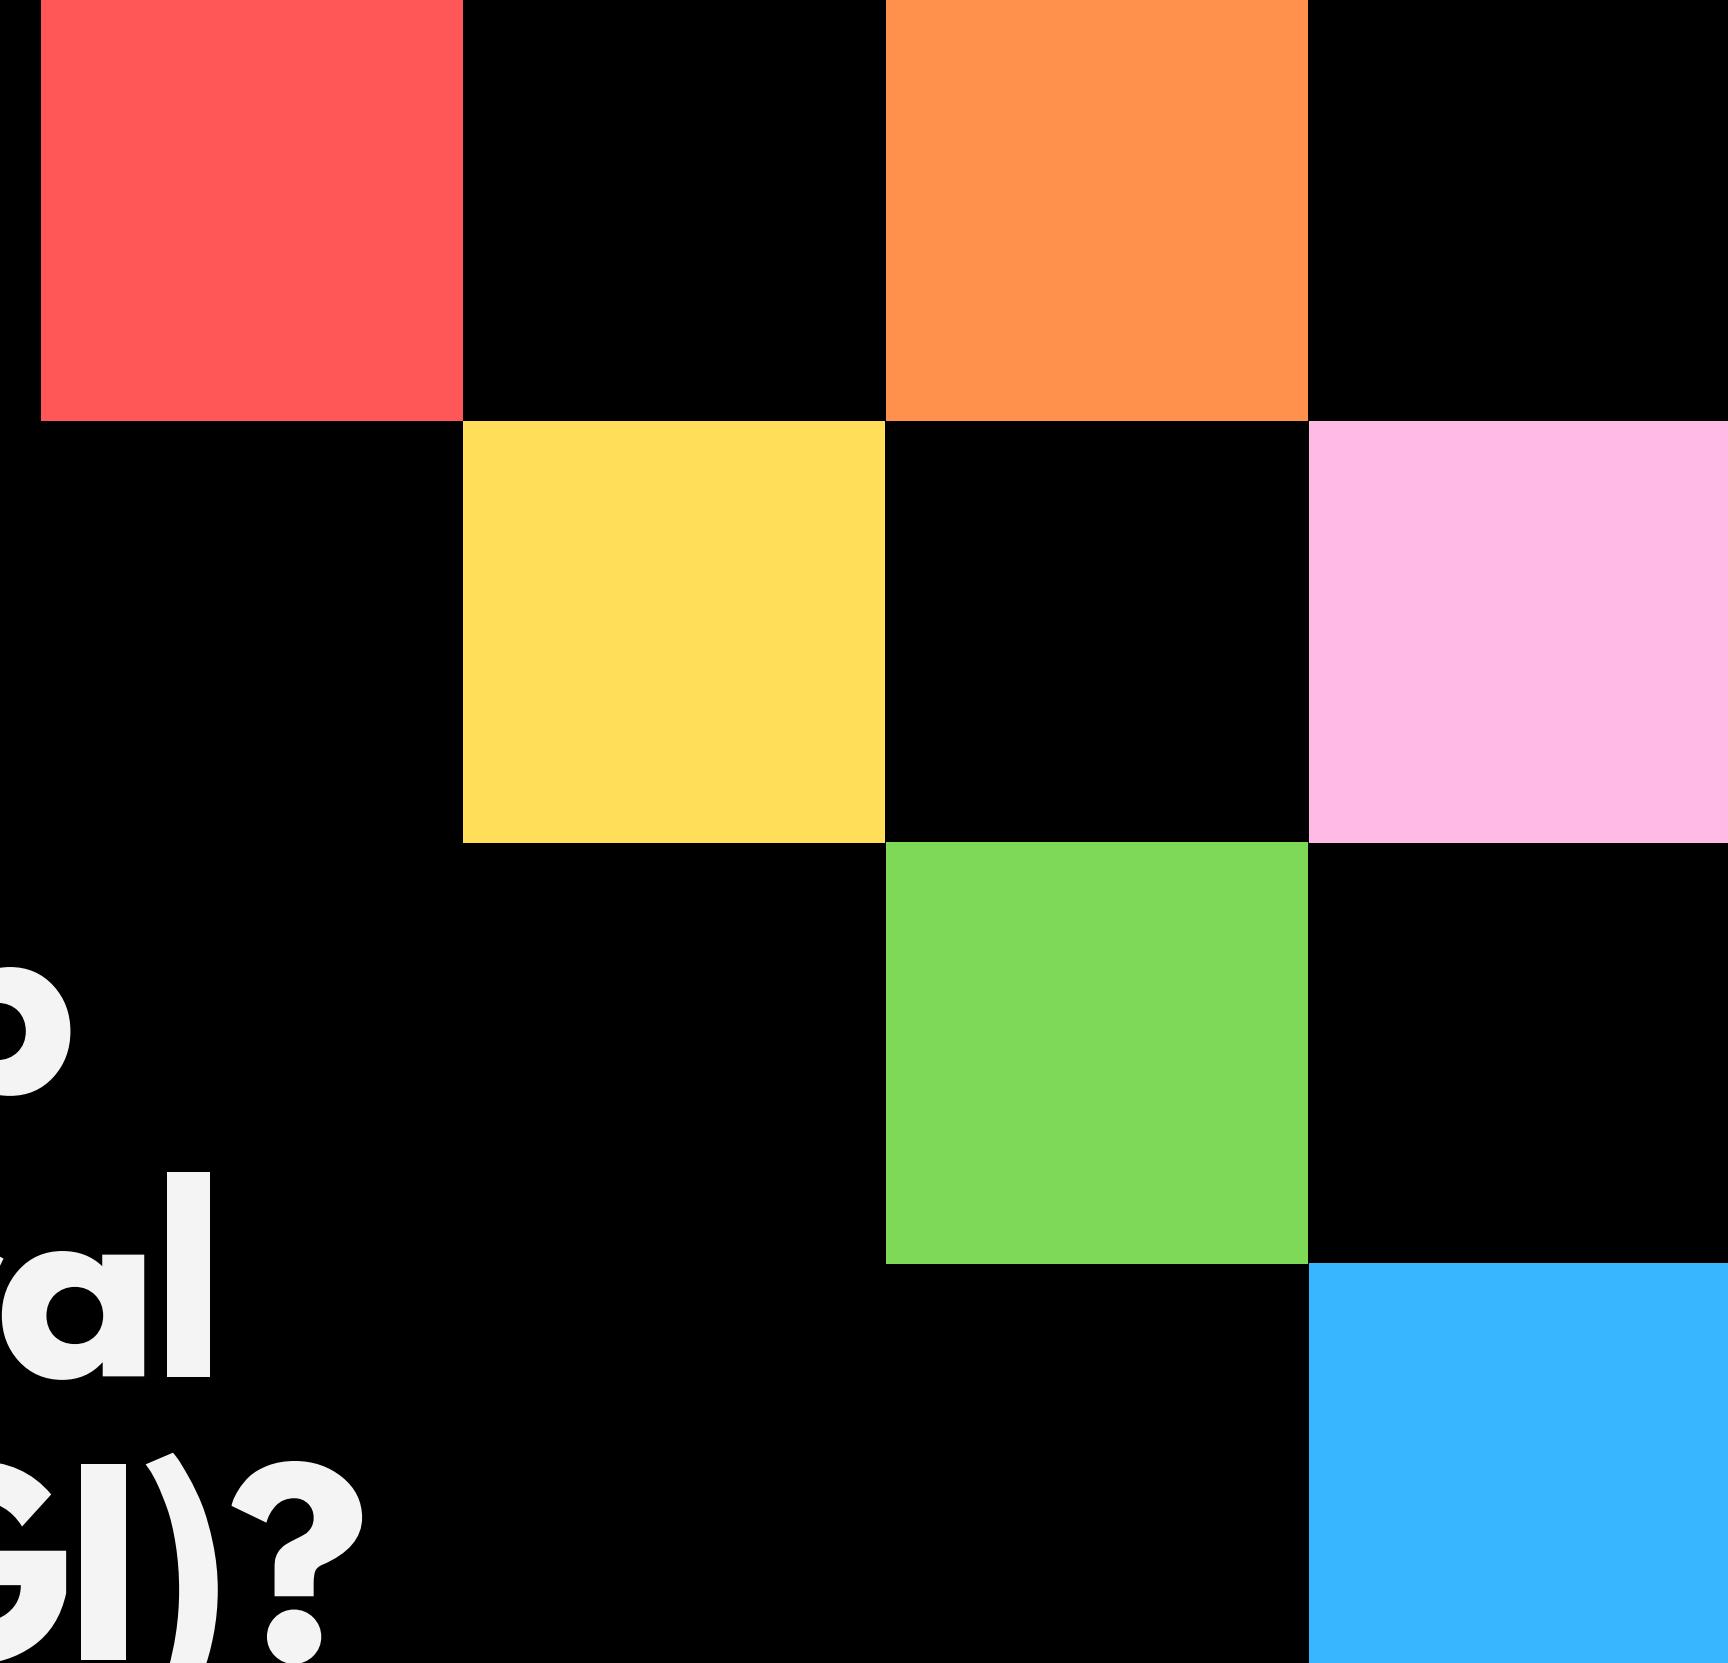


and 1000+ unique problems!



Problem 1

**How to increase
number of unique
problems solved?**



Problem 2

**How to develop
Artificial General
Intelligence (AGI)?**



**How to do this as a
student?**

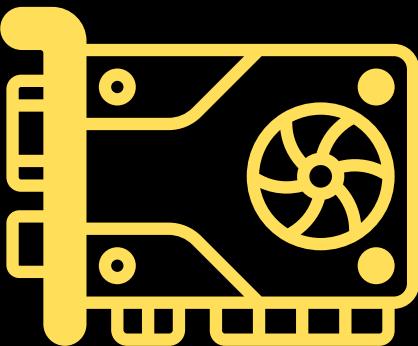
Solution Cost

\$0

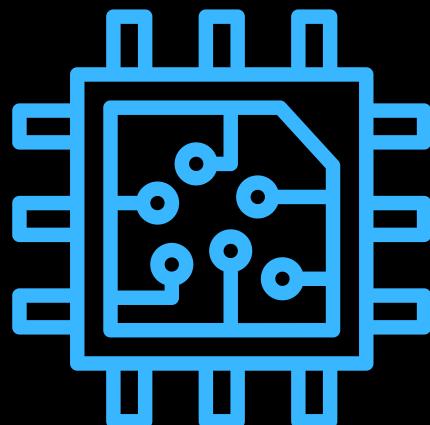
System Configuration



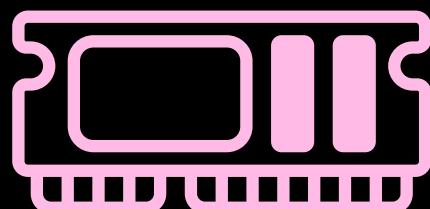
Mizzou HellBender



Node type: GPU (H100)



Number of CPUs: 24

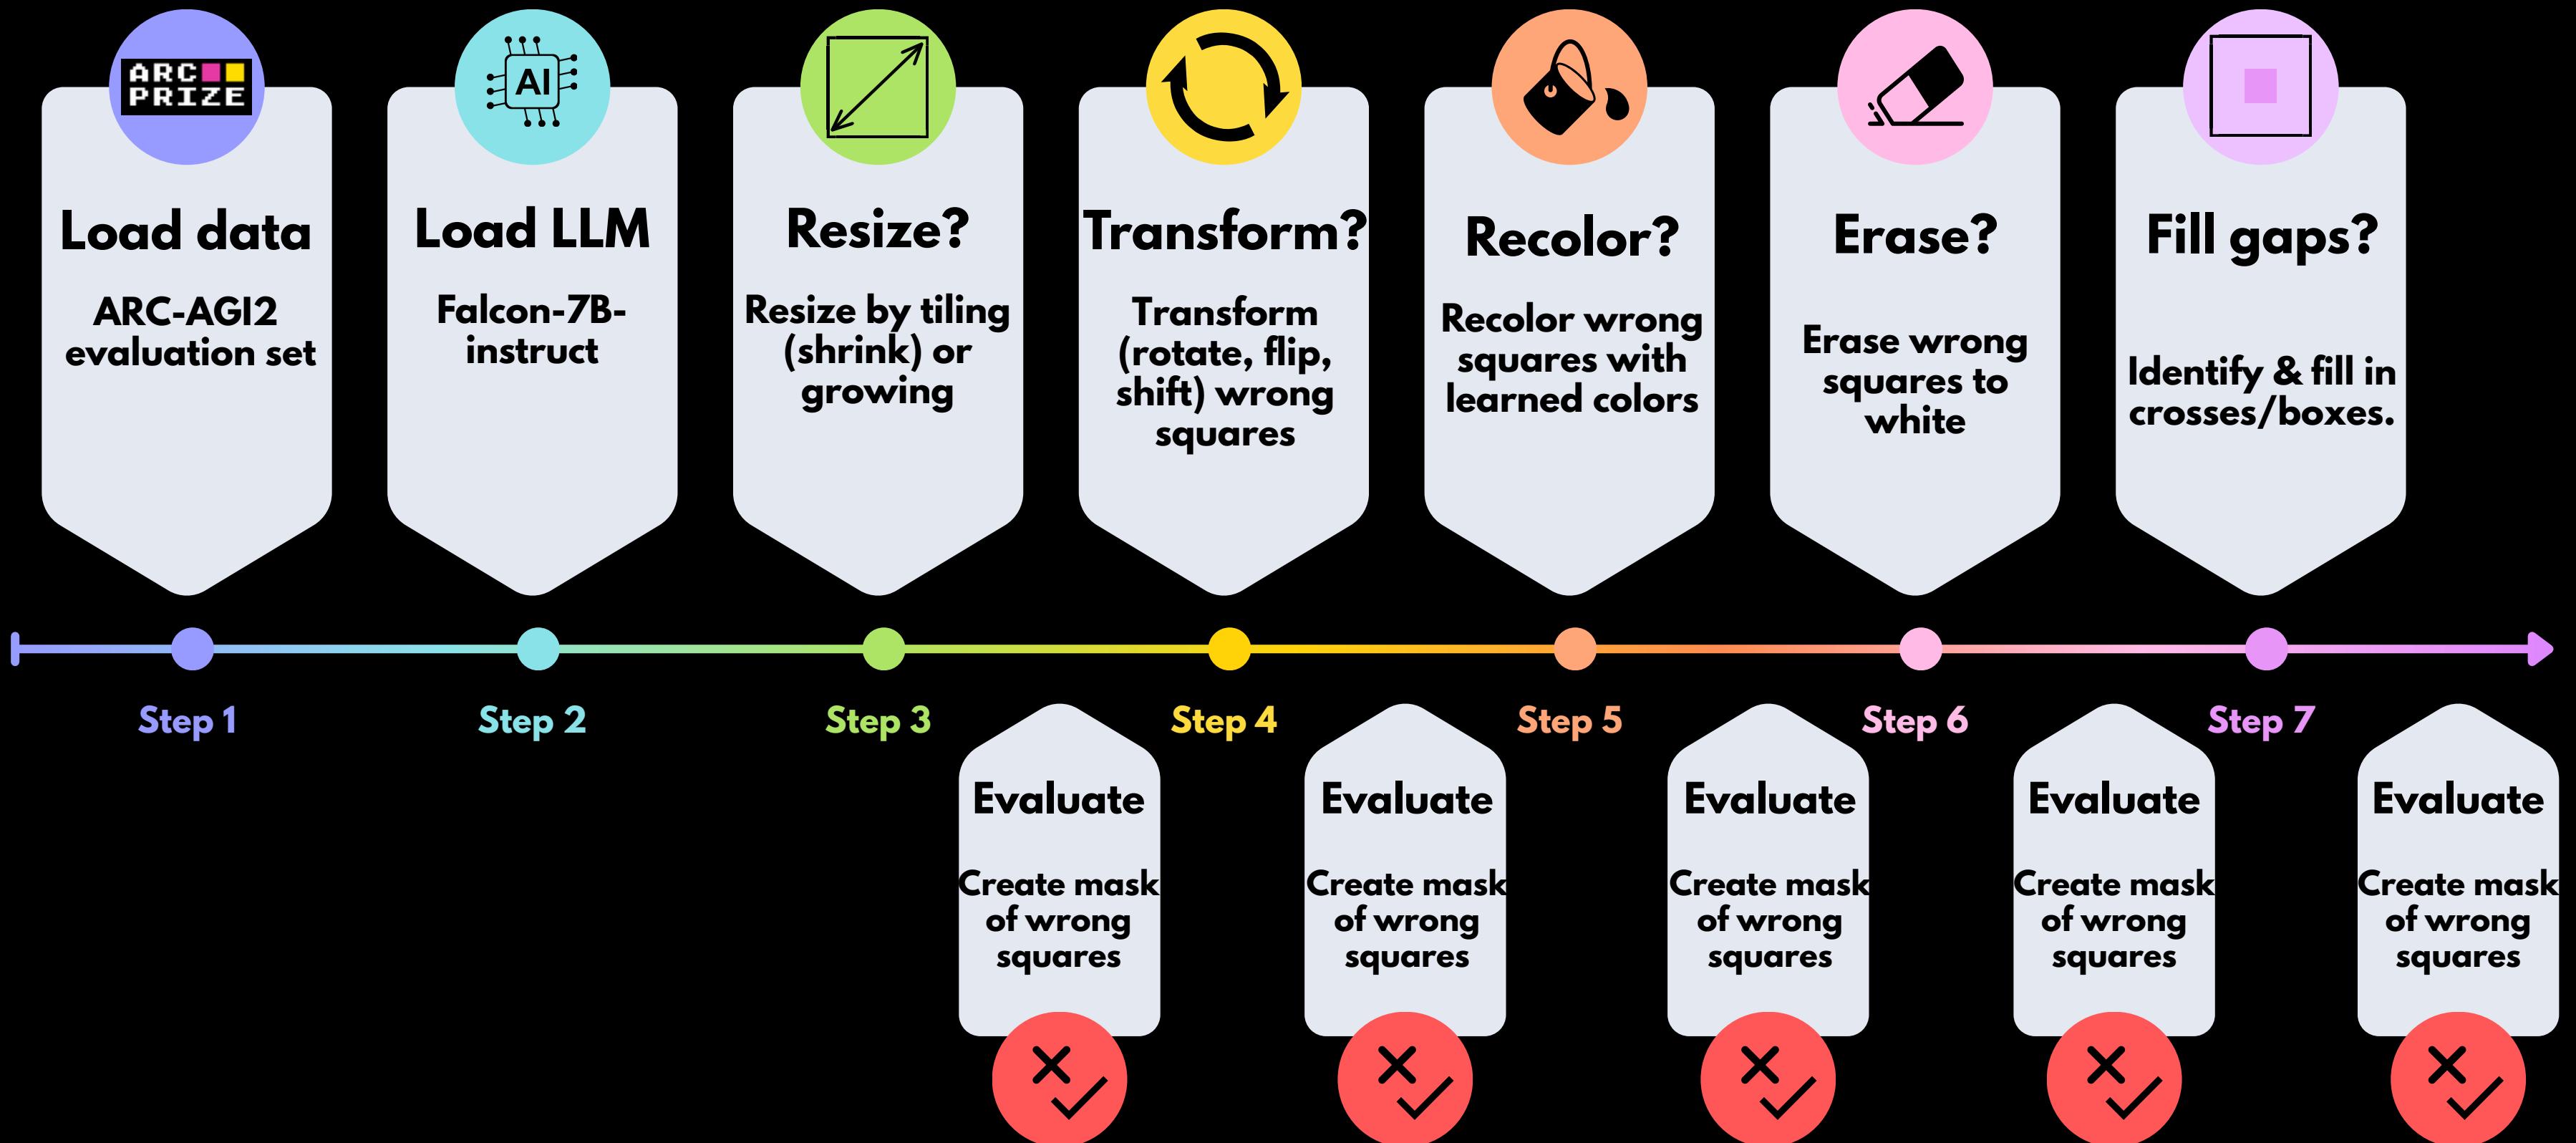


Memory: 32 GB

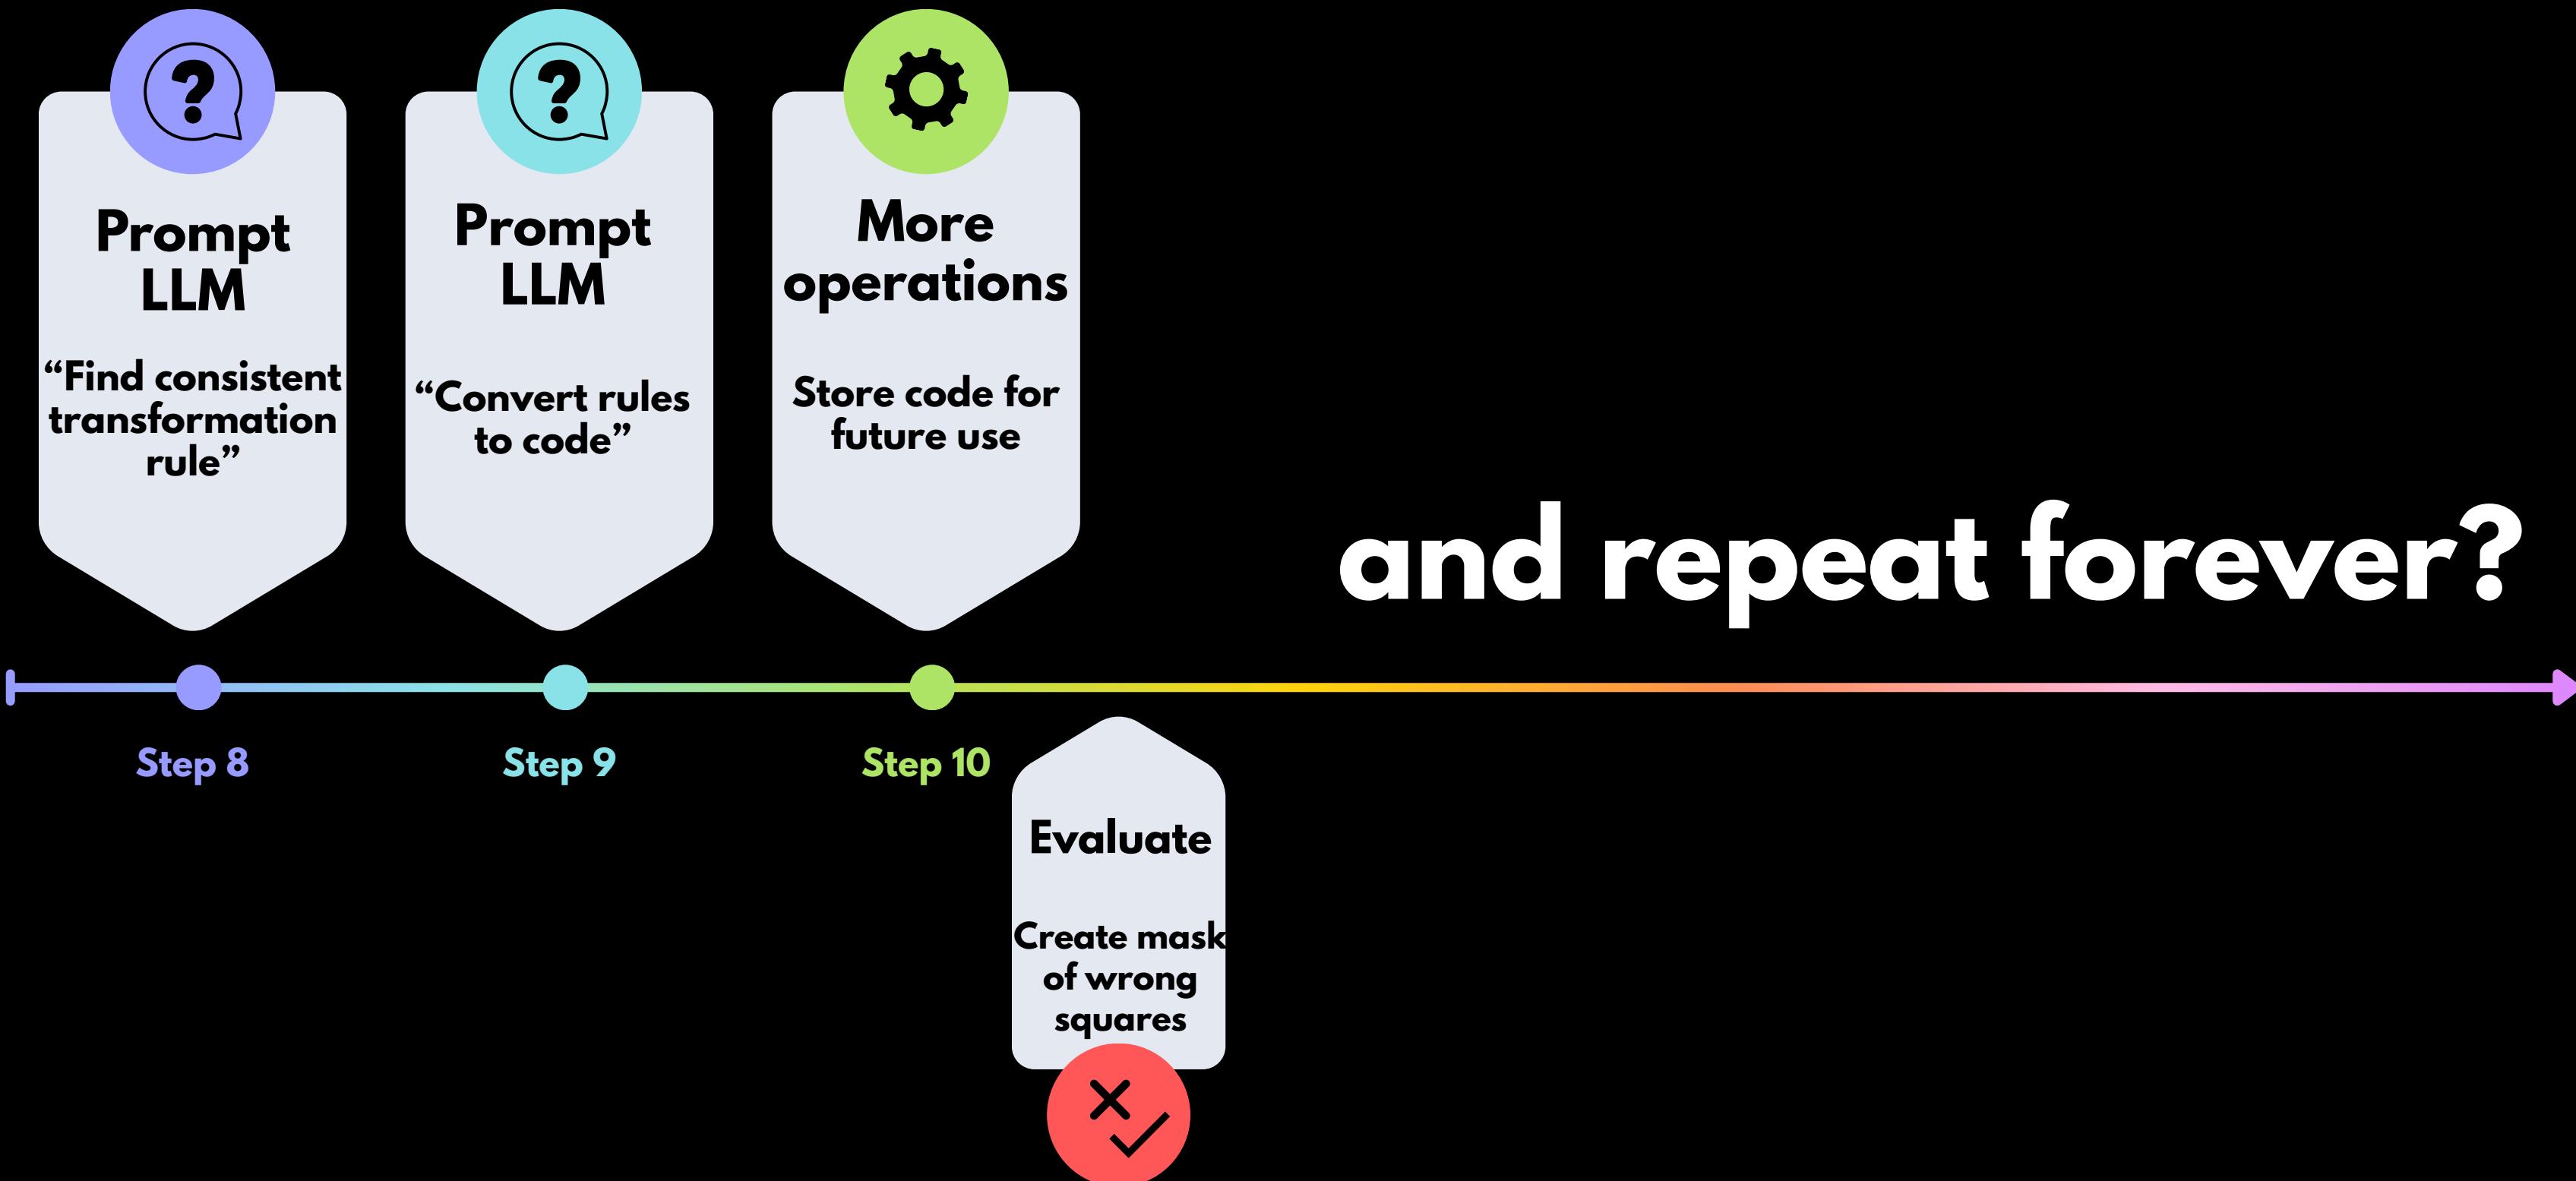
Question 1

**How to increase
number of unique
problems solved?**

Methodology



Methodology



> LLM setup

Falcon-7B (H100 Optimized)

↳ 1 cell hidden



> Operations setup

Load Resize, Transforms, Recolor, Erase, FillGaps, etc.

▶ ↳ 7 cells hidden

> Execute methodology

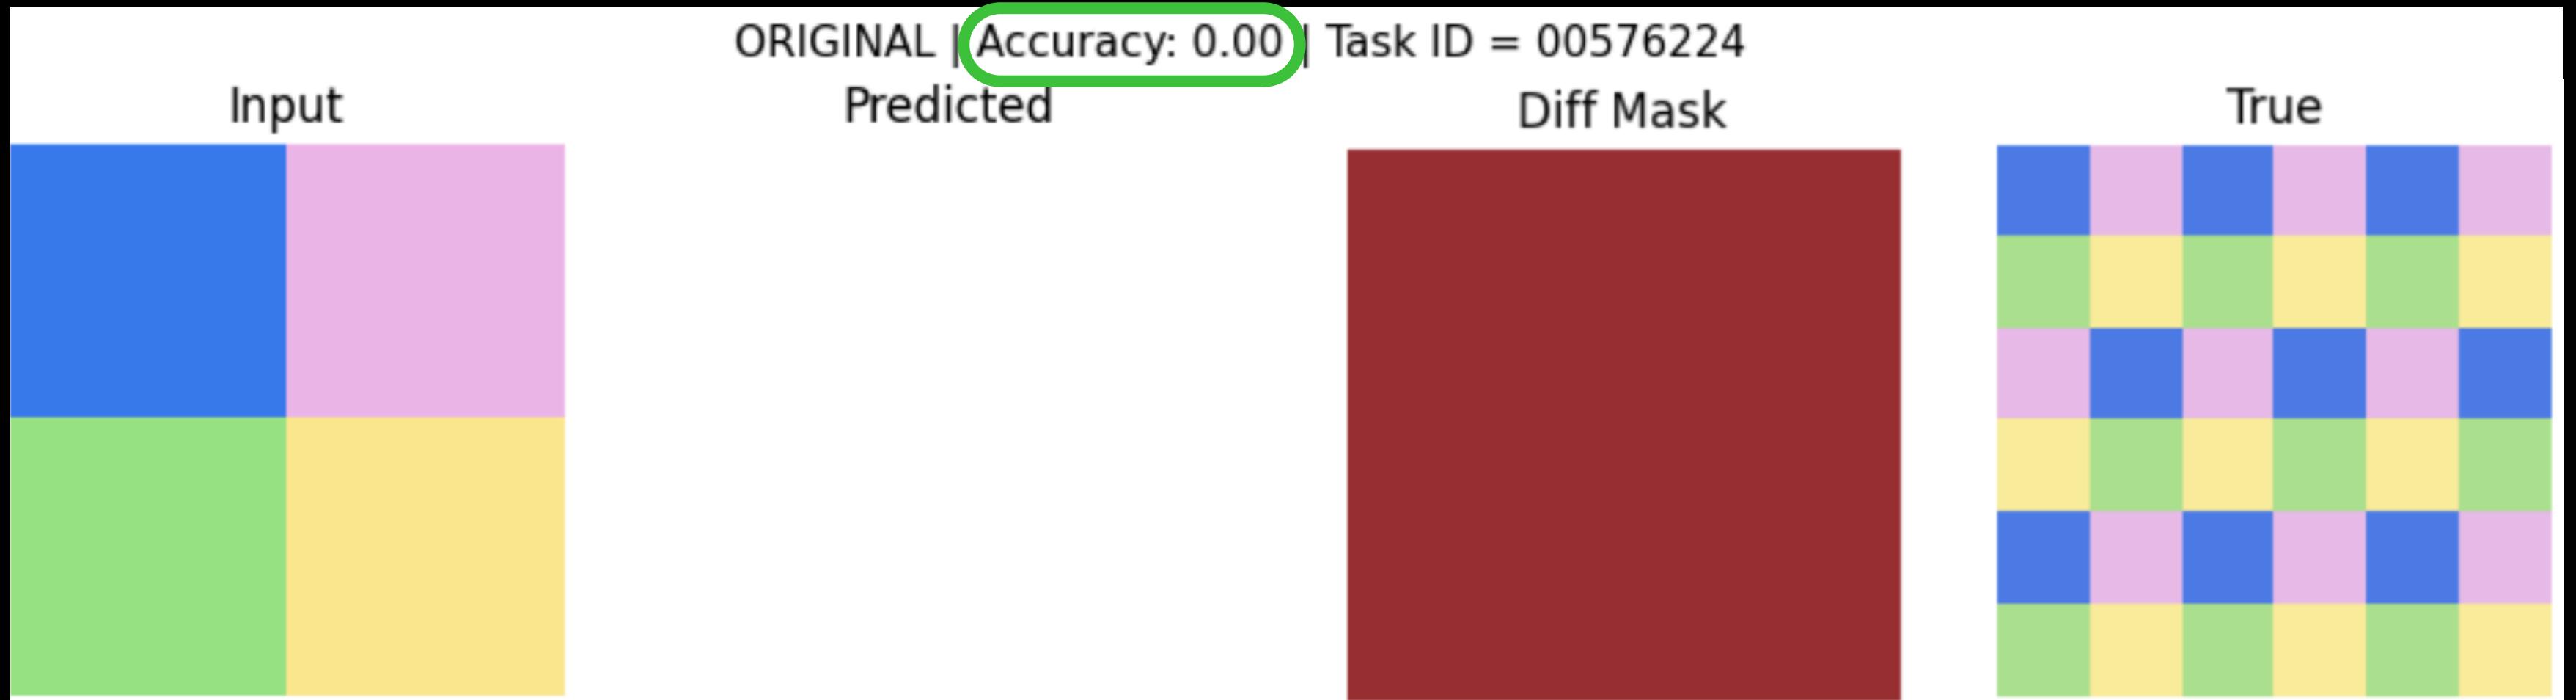
[11]

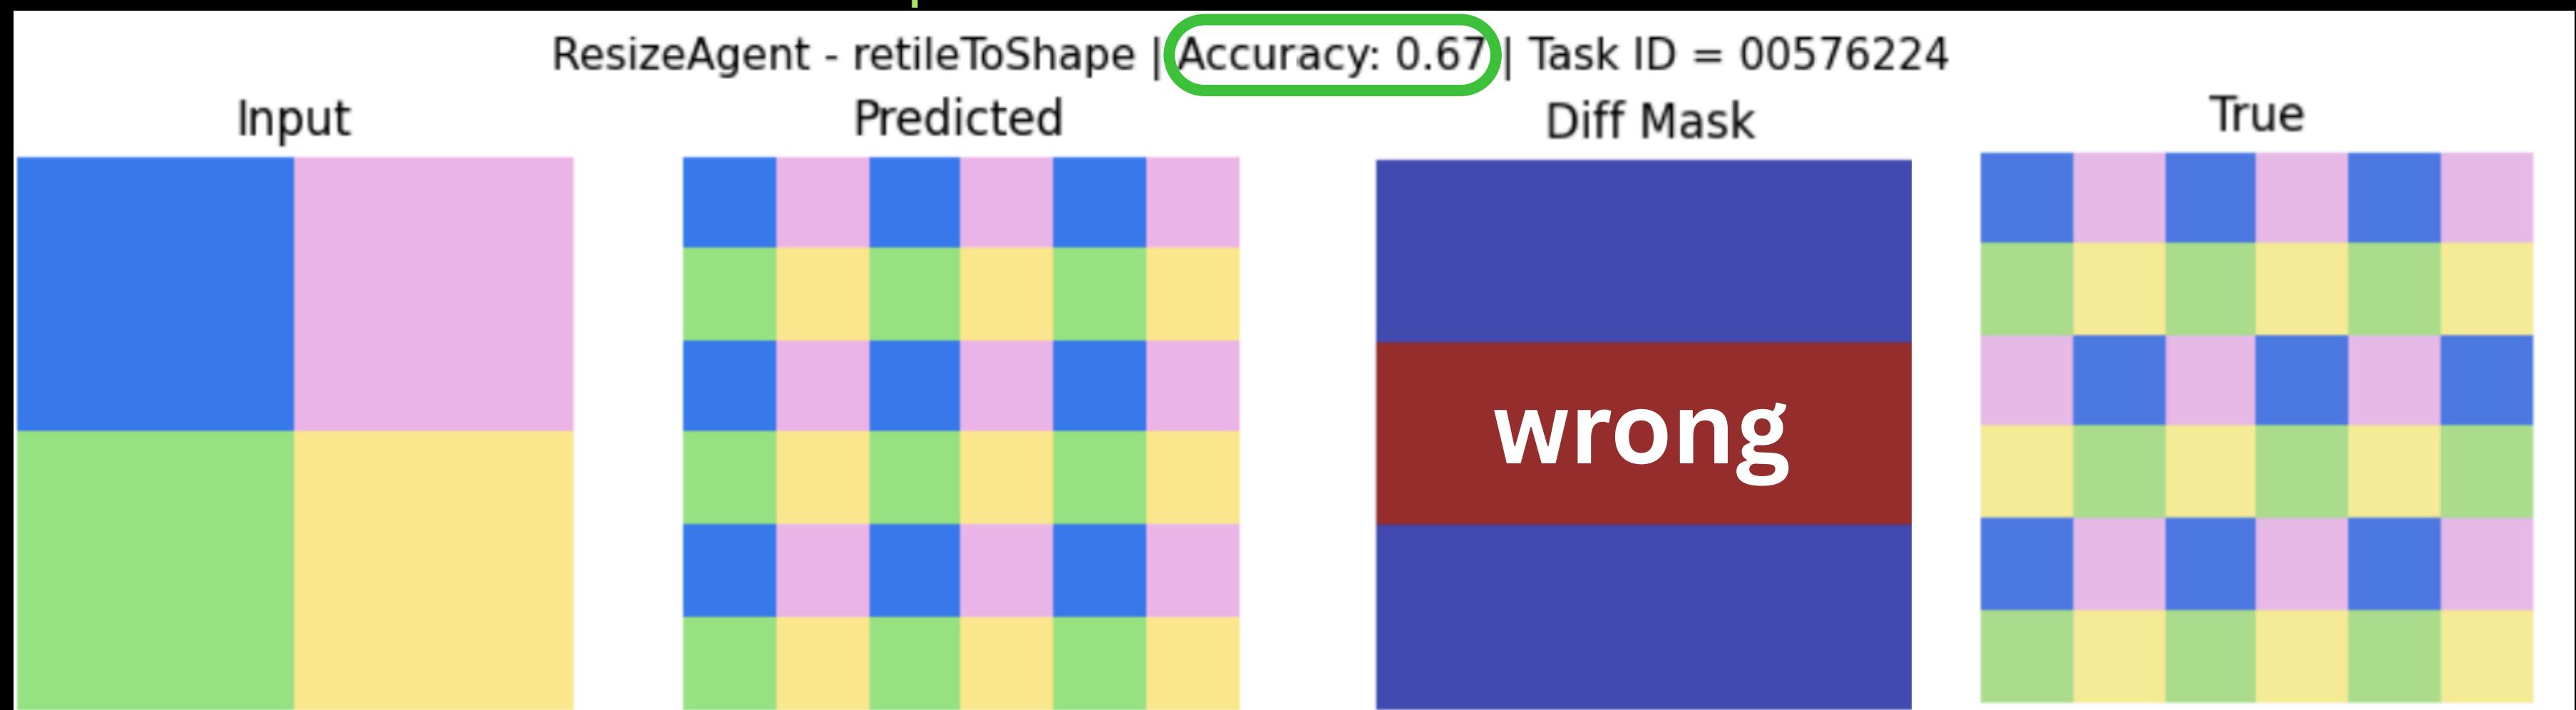
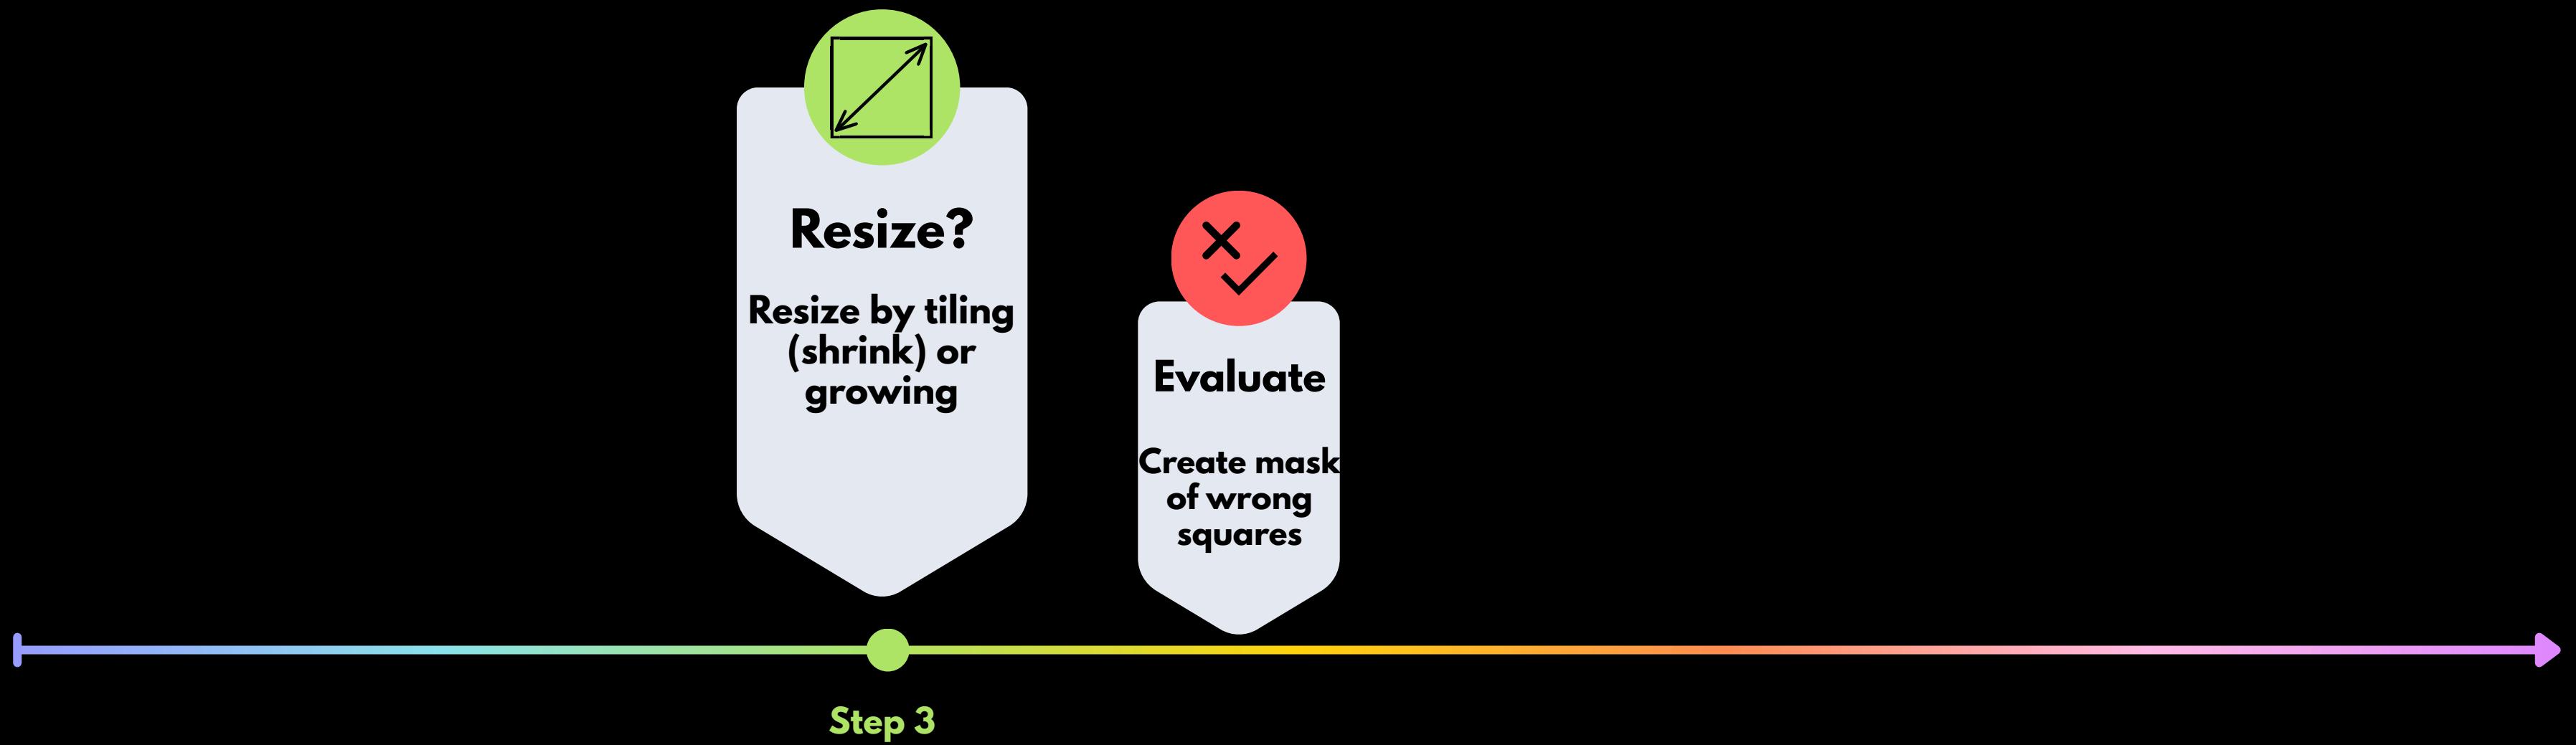
! 57s

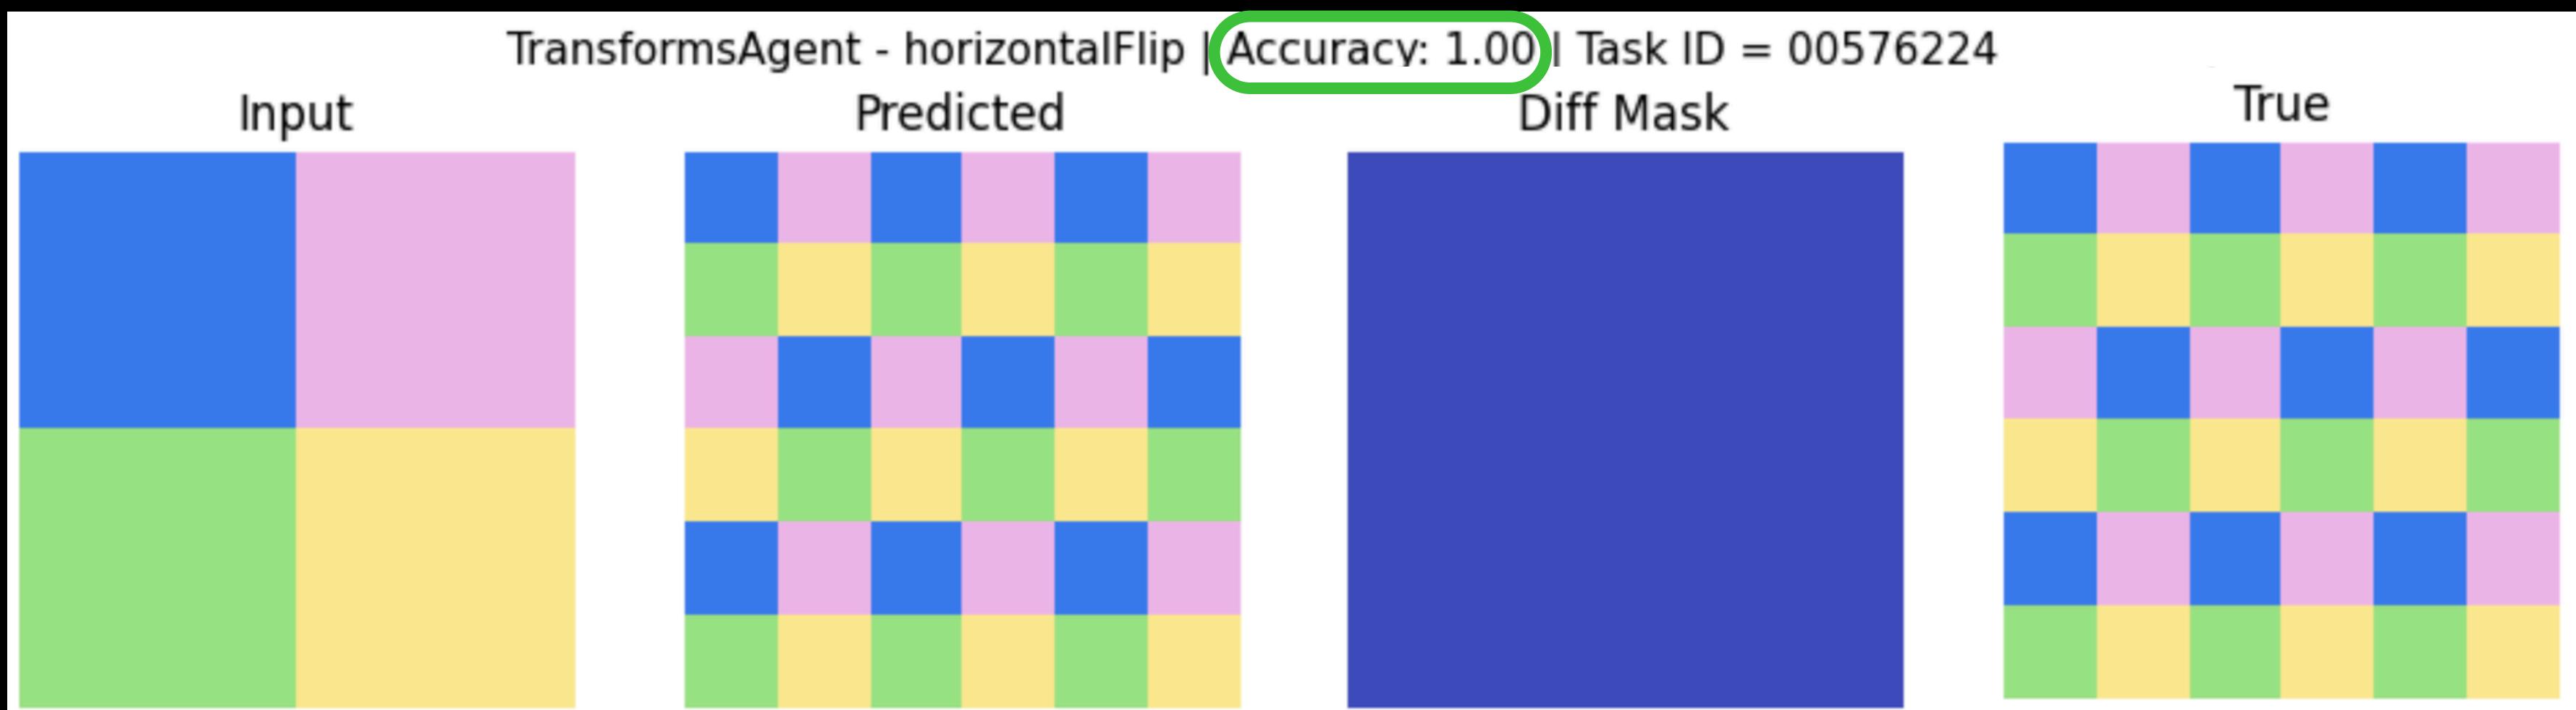
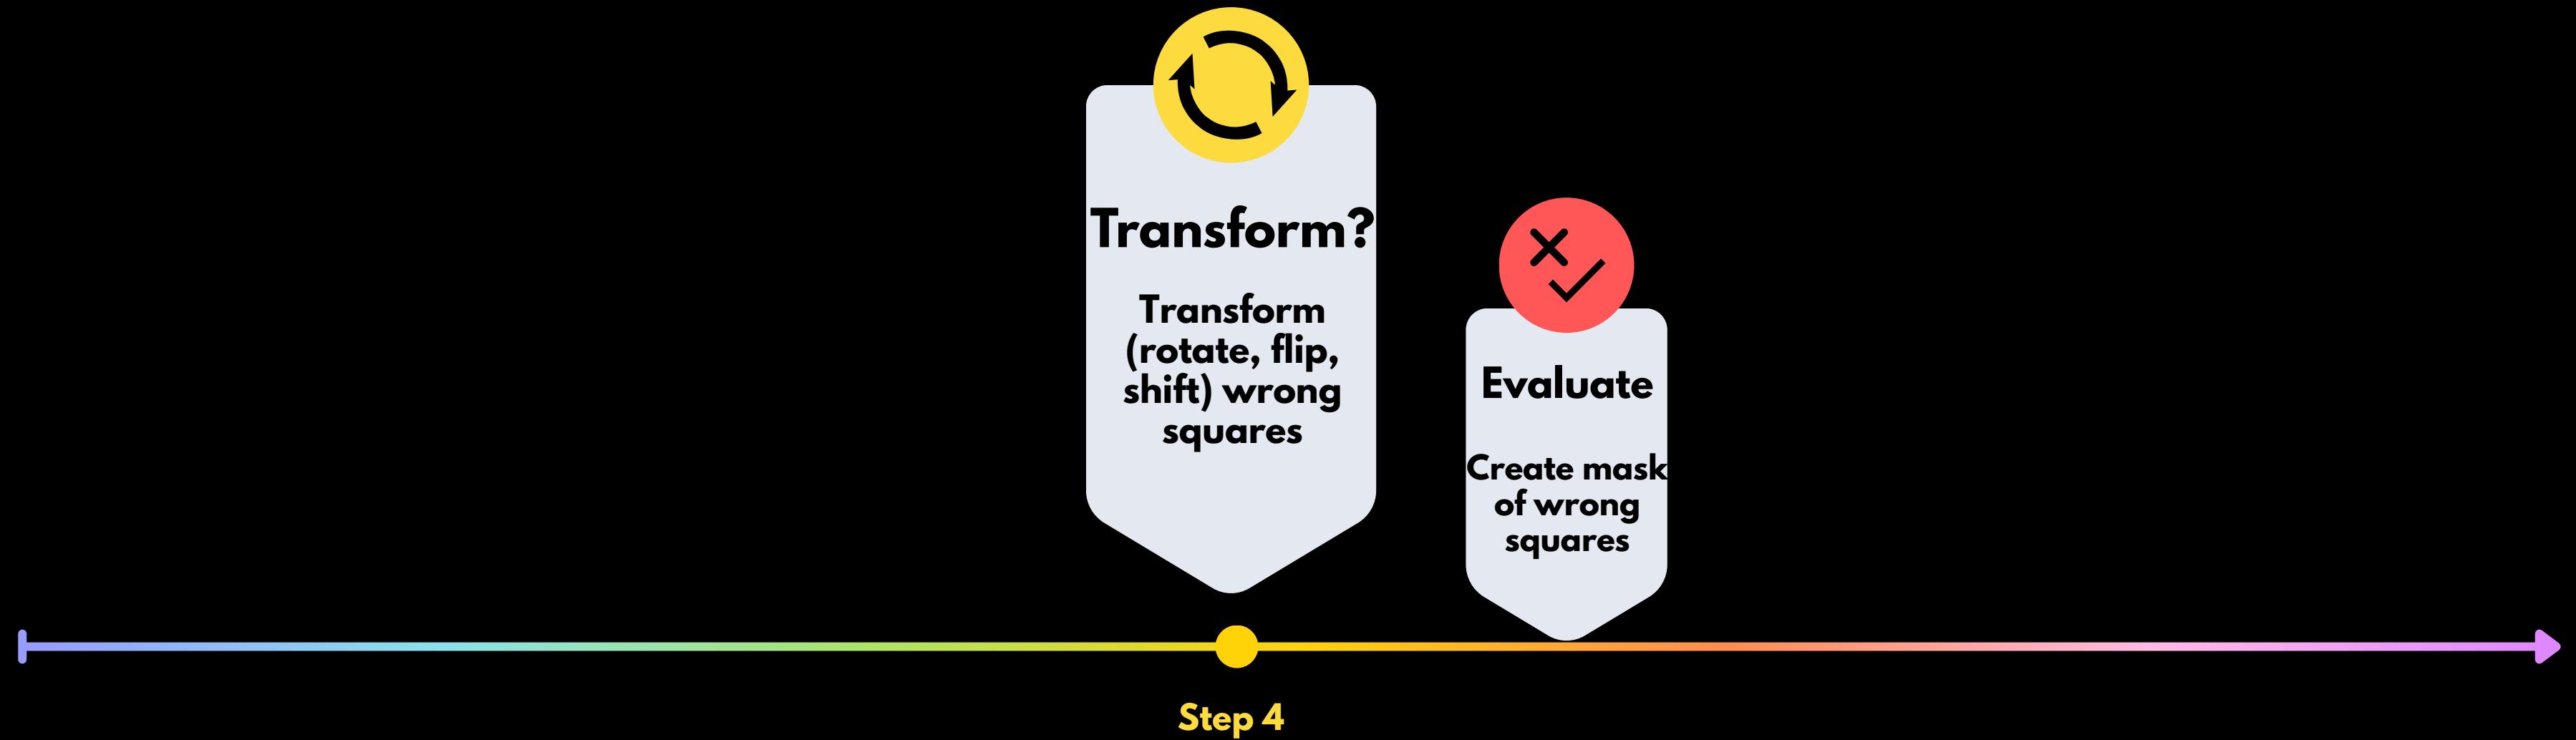
```
1 # Lets try first few examples
2 start = 0
3 stop = 6
4 x= 0
5 test_one = False
6 for id in train_ids[start:stop]:
7     # Fetch test from train
8         test_arr, label_arr = train_data[id]["test"][0]["input"], train_label[id][0]
9         test_arr, label_arr = convert_2D_numpy(test_arr), convert_2D_numpy(label_arr)
10
11 # Define Agents
```

~1 min to solve first 6 unique problems

[Play slide 14 video](#)







Results on evaluation problems

Pixel Correctness (Fine-Grained Accuracy)

91.23%

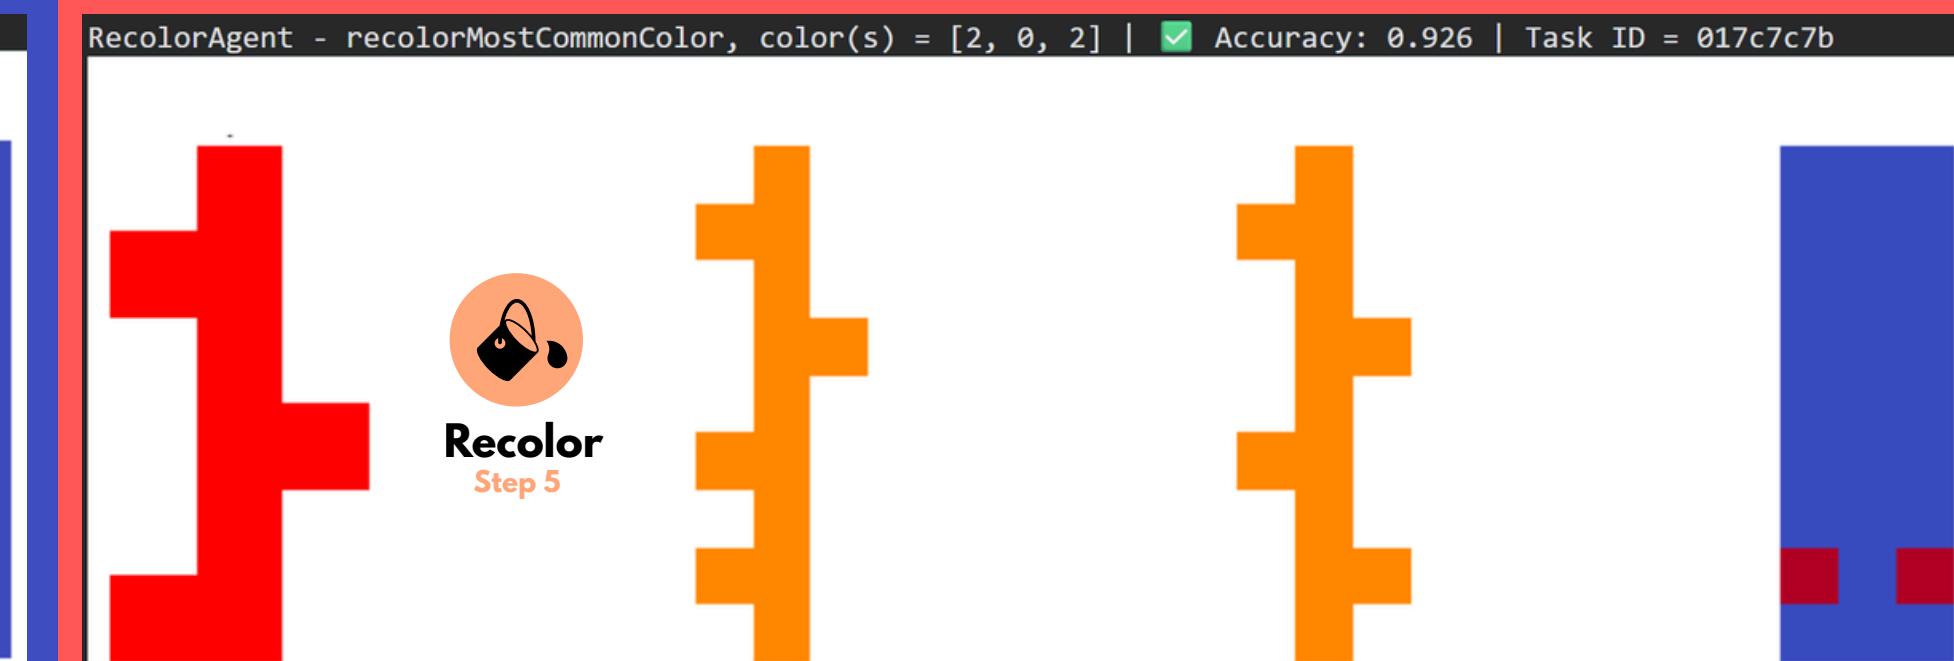
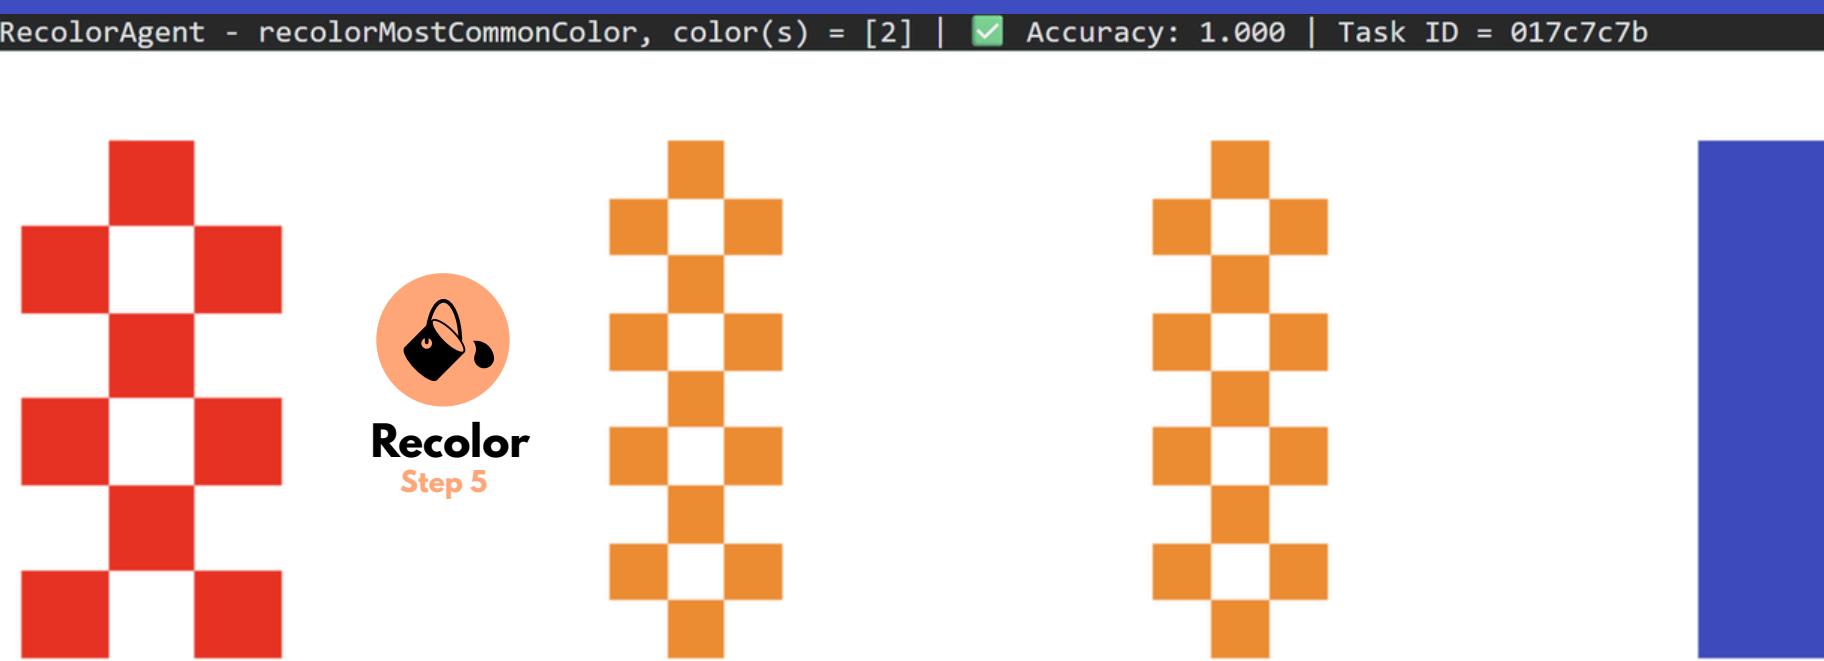
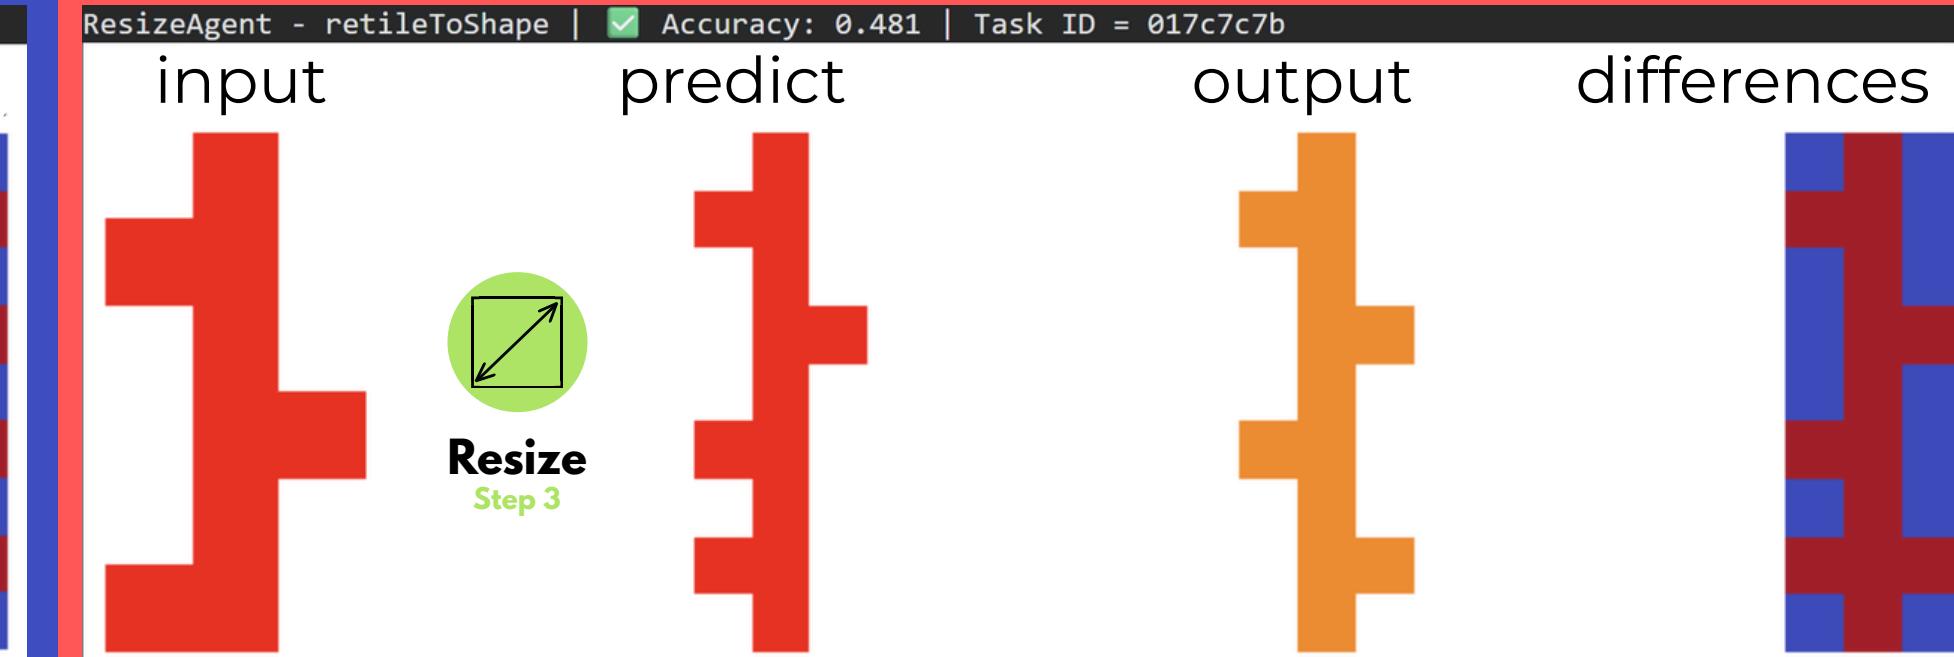
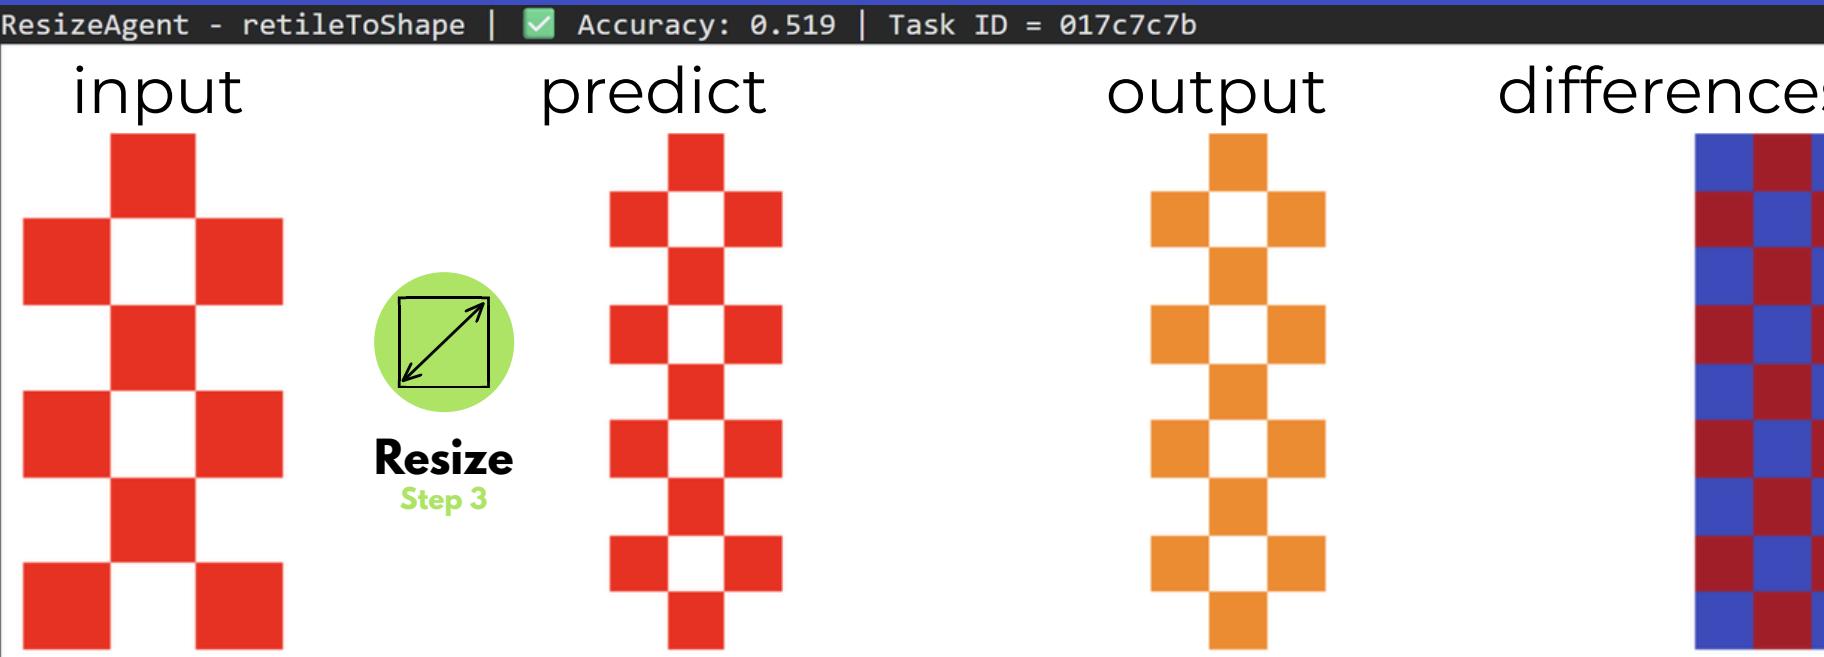
across each example

Correct / Incorrect (Task-Level Accuracy)

10.59%

120 problems total

Even within same problem, operations can vary



Correct



Incorrect



Step 8

Step 9



Prompt LLM

“Find consistent transformation rule”



Prompt LLM

“Convert rules to code”

prompt = “Example1: Agents can't solve task. Best accuracy=0.93 using operations [<{'RecolorAgent': ['recolorMostCommonColor', color(s) = [2, 0, 2]}]}].

best_predict_arr=[[0 2 0]...]

input_arr=[[0 1 0]...]

true_arr=[[0 2 0]...]

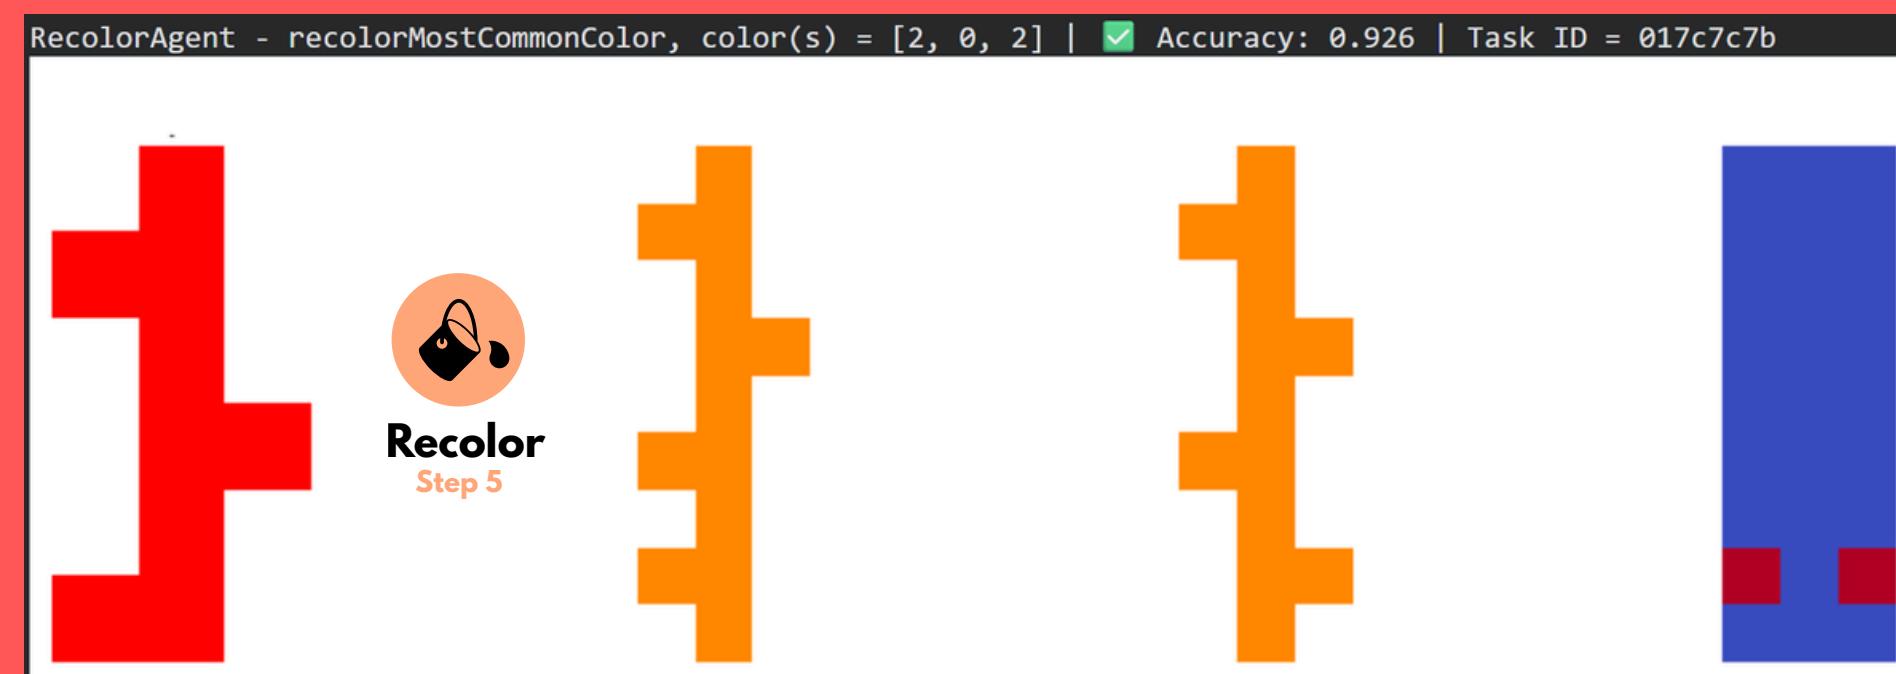
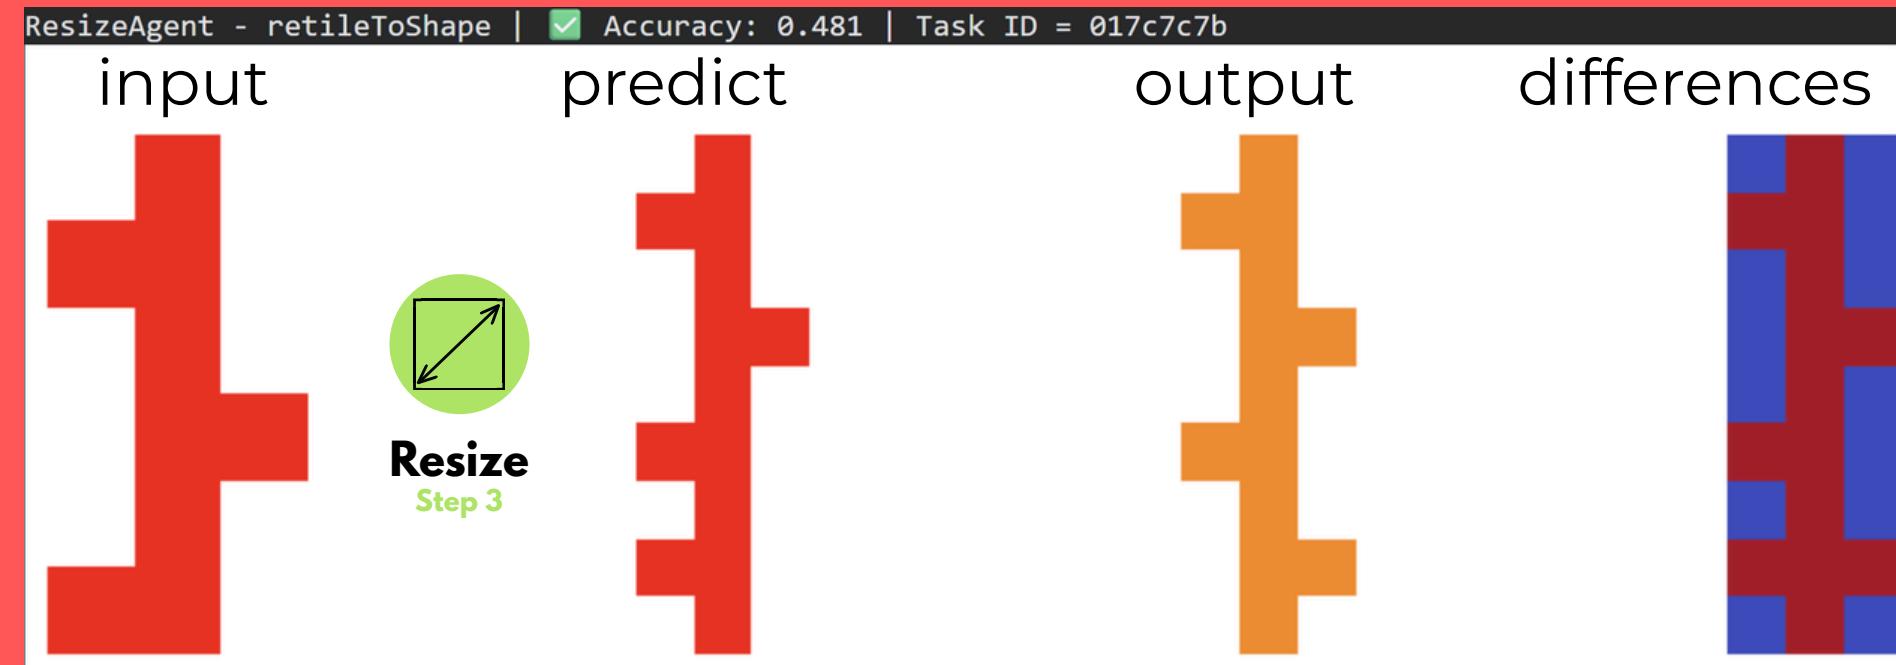
differences=[[False False False]... [True False True]...]

1. Identify what is still different

2. List the reasoning behind input_arr -> true_arr?

3. Generalize reasoning so it'll work for other similar examples

4. Generate python code that would solve this task”



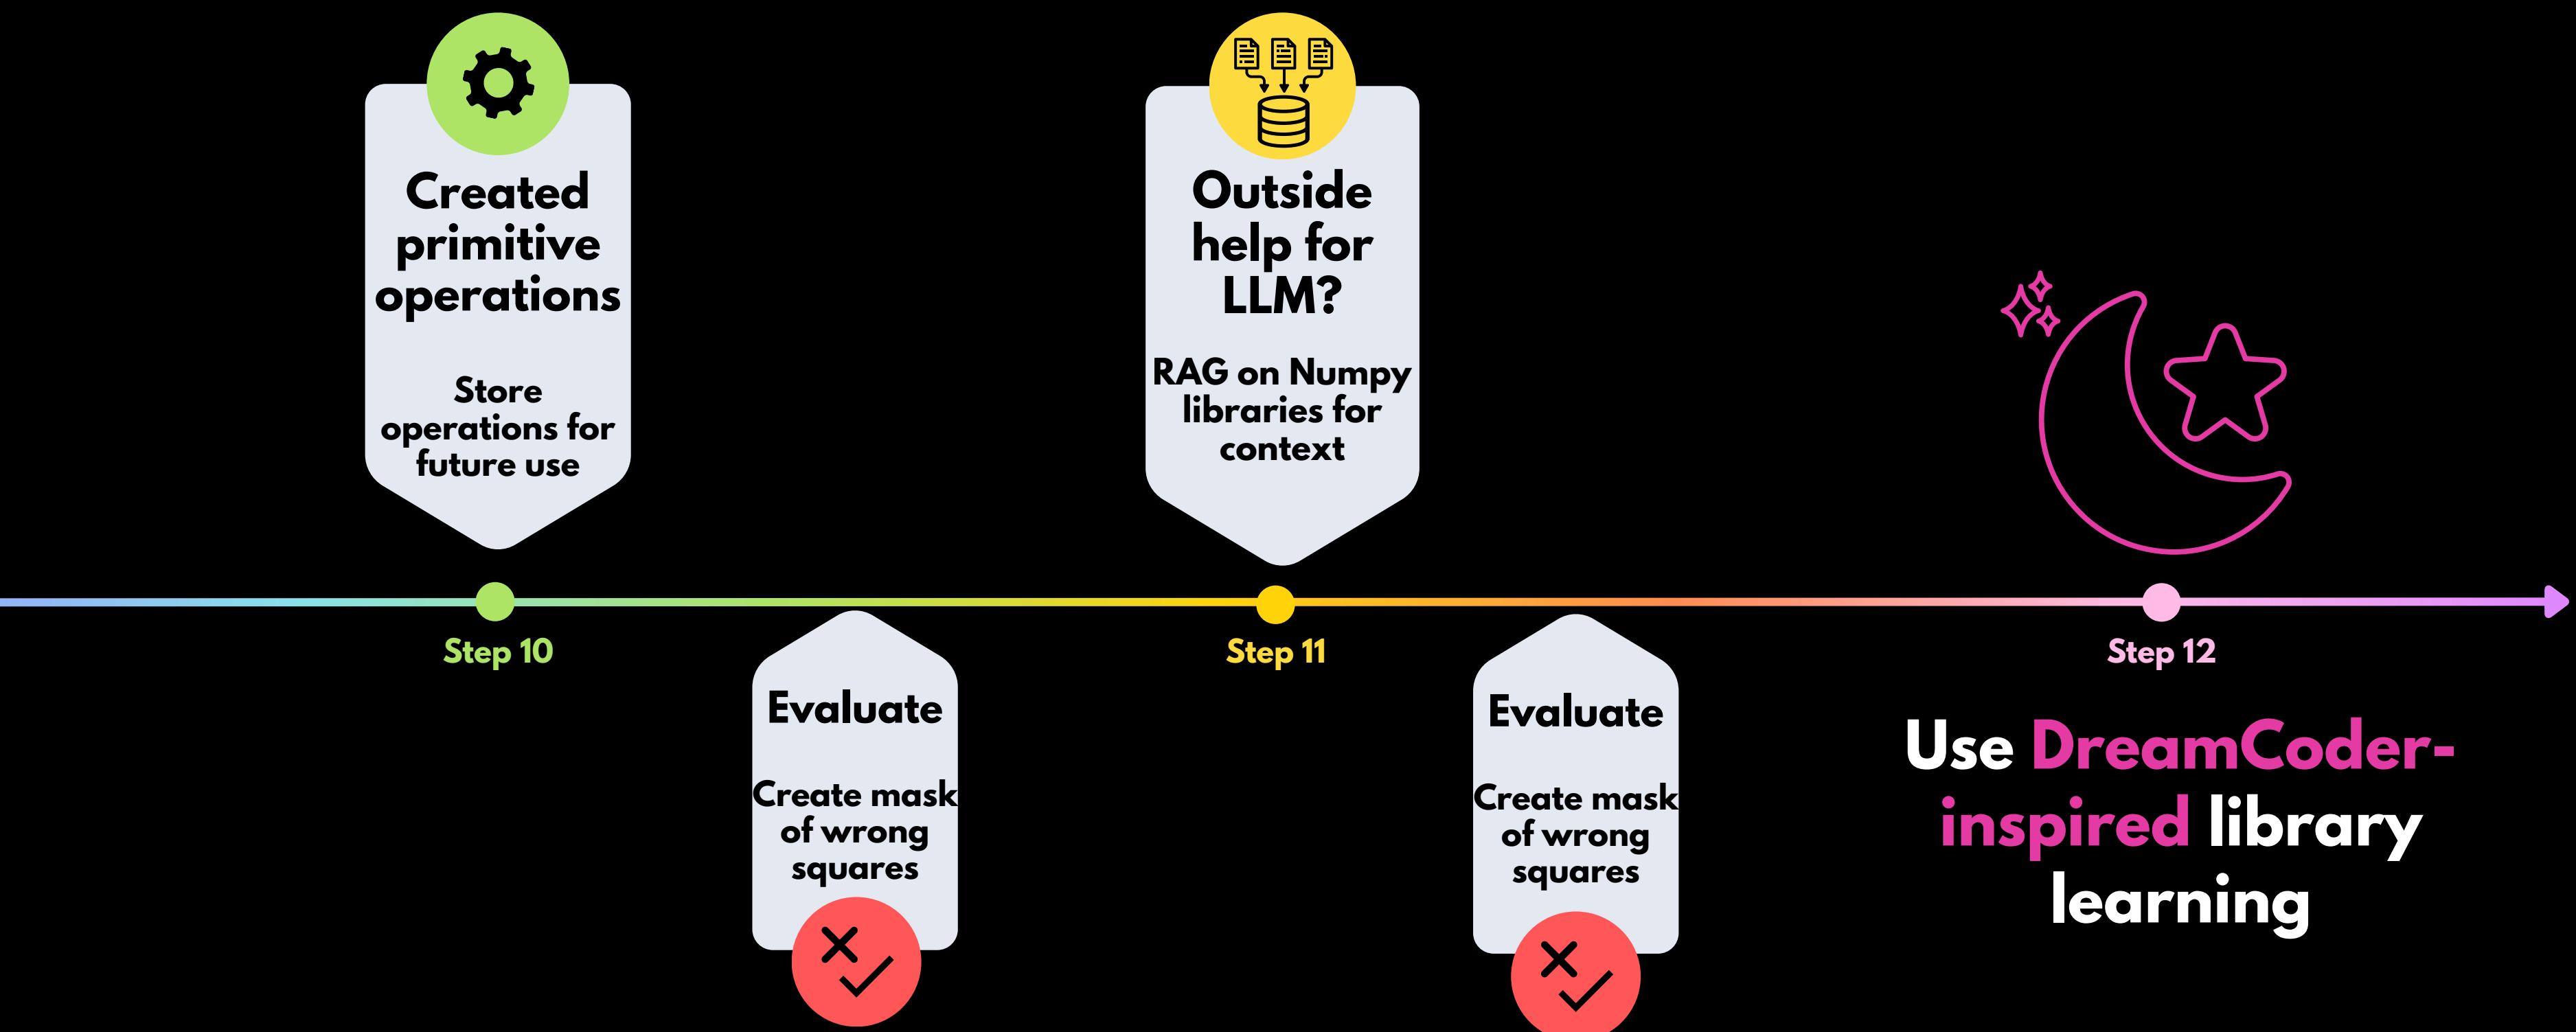
Incorrect



Question 2

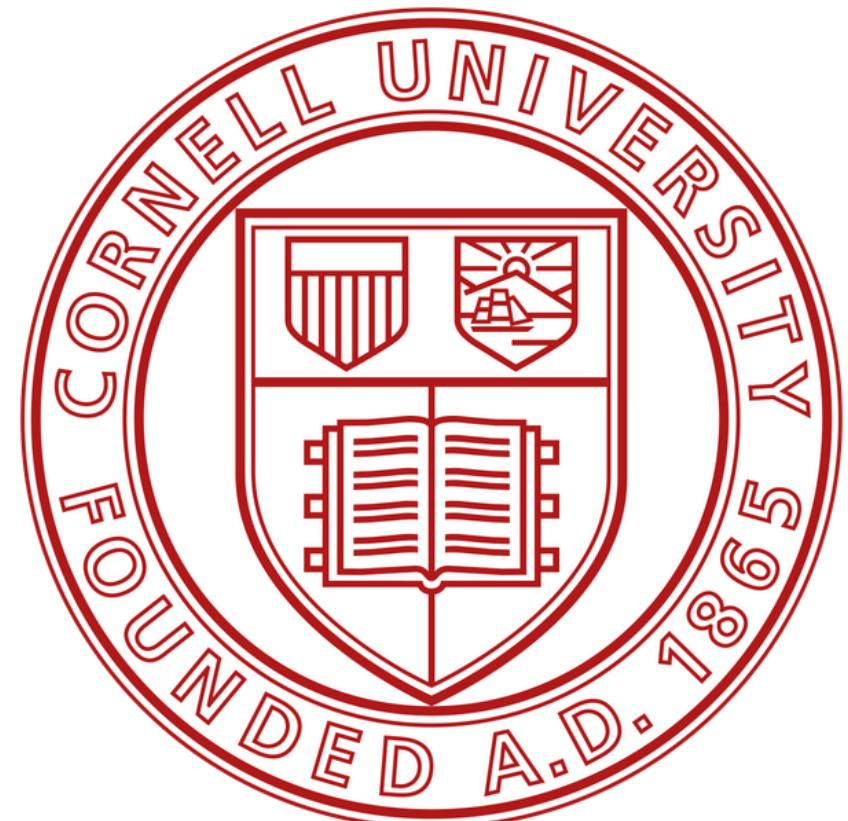
**How to develop
Artificial General
Intelligence (AGI)?**

Methodology for AGI



Motivation of adapting DreamCoder

- Problems are combinatorial & symbolic (colors, shapes, symmetries, object operations)
- We want compositional generalization: reuse concepts like rotate, crop, repeat across puzzles



Motivation of adapting DreamCoder

Why DreamCoder-style approach?

- System should learn building blocks, not just simply memorize per-task tricks
- A standard neural network outputs a grid directly and is hard to interpret

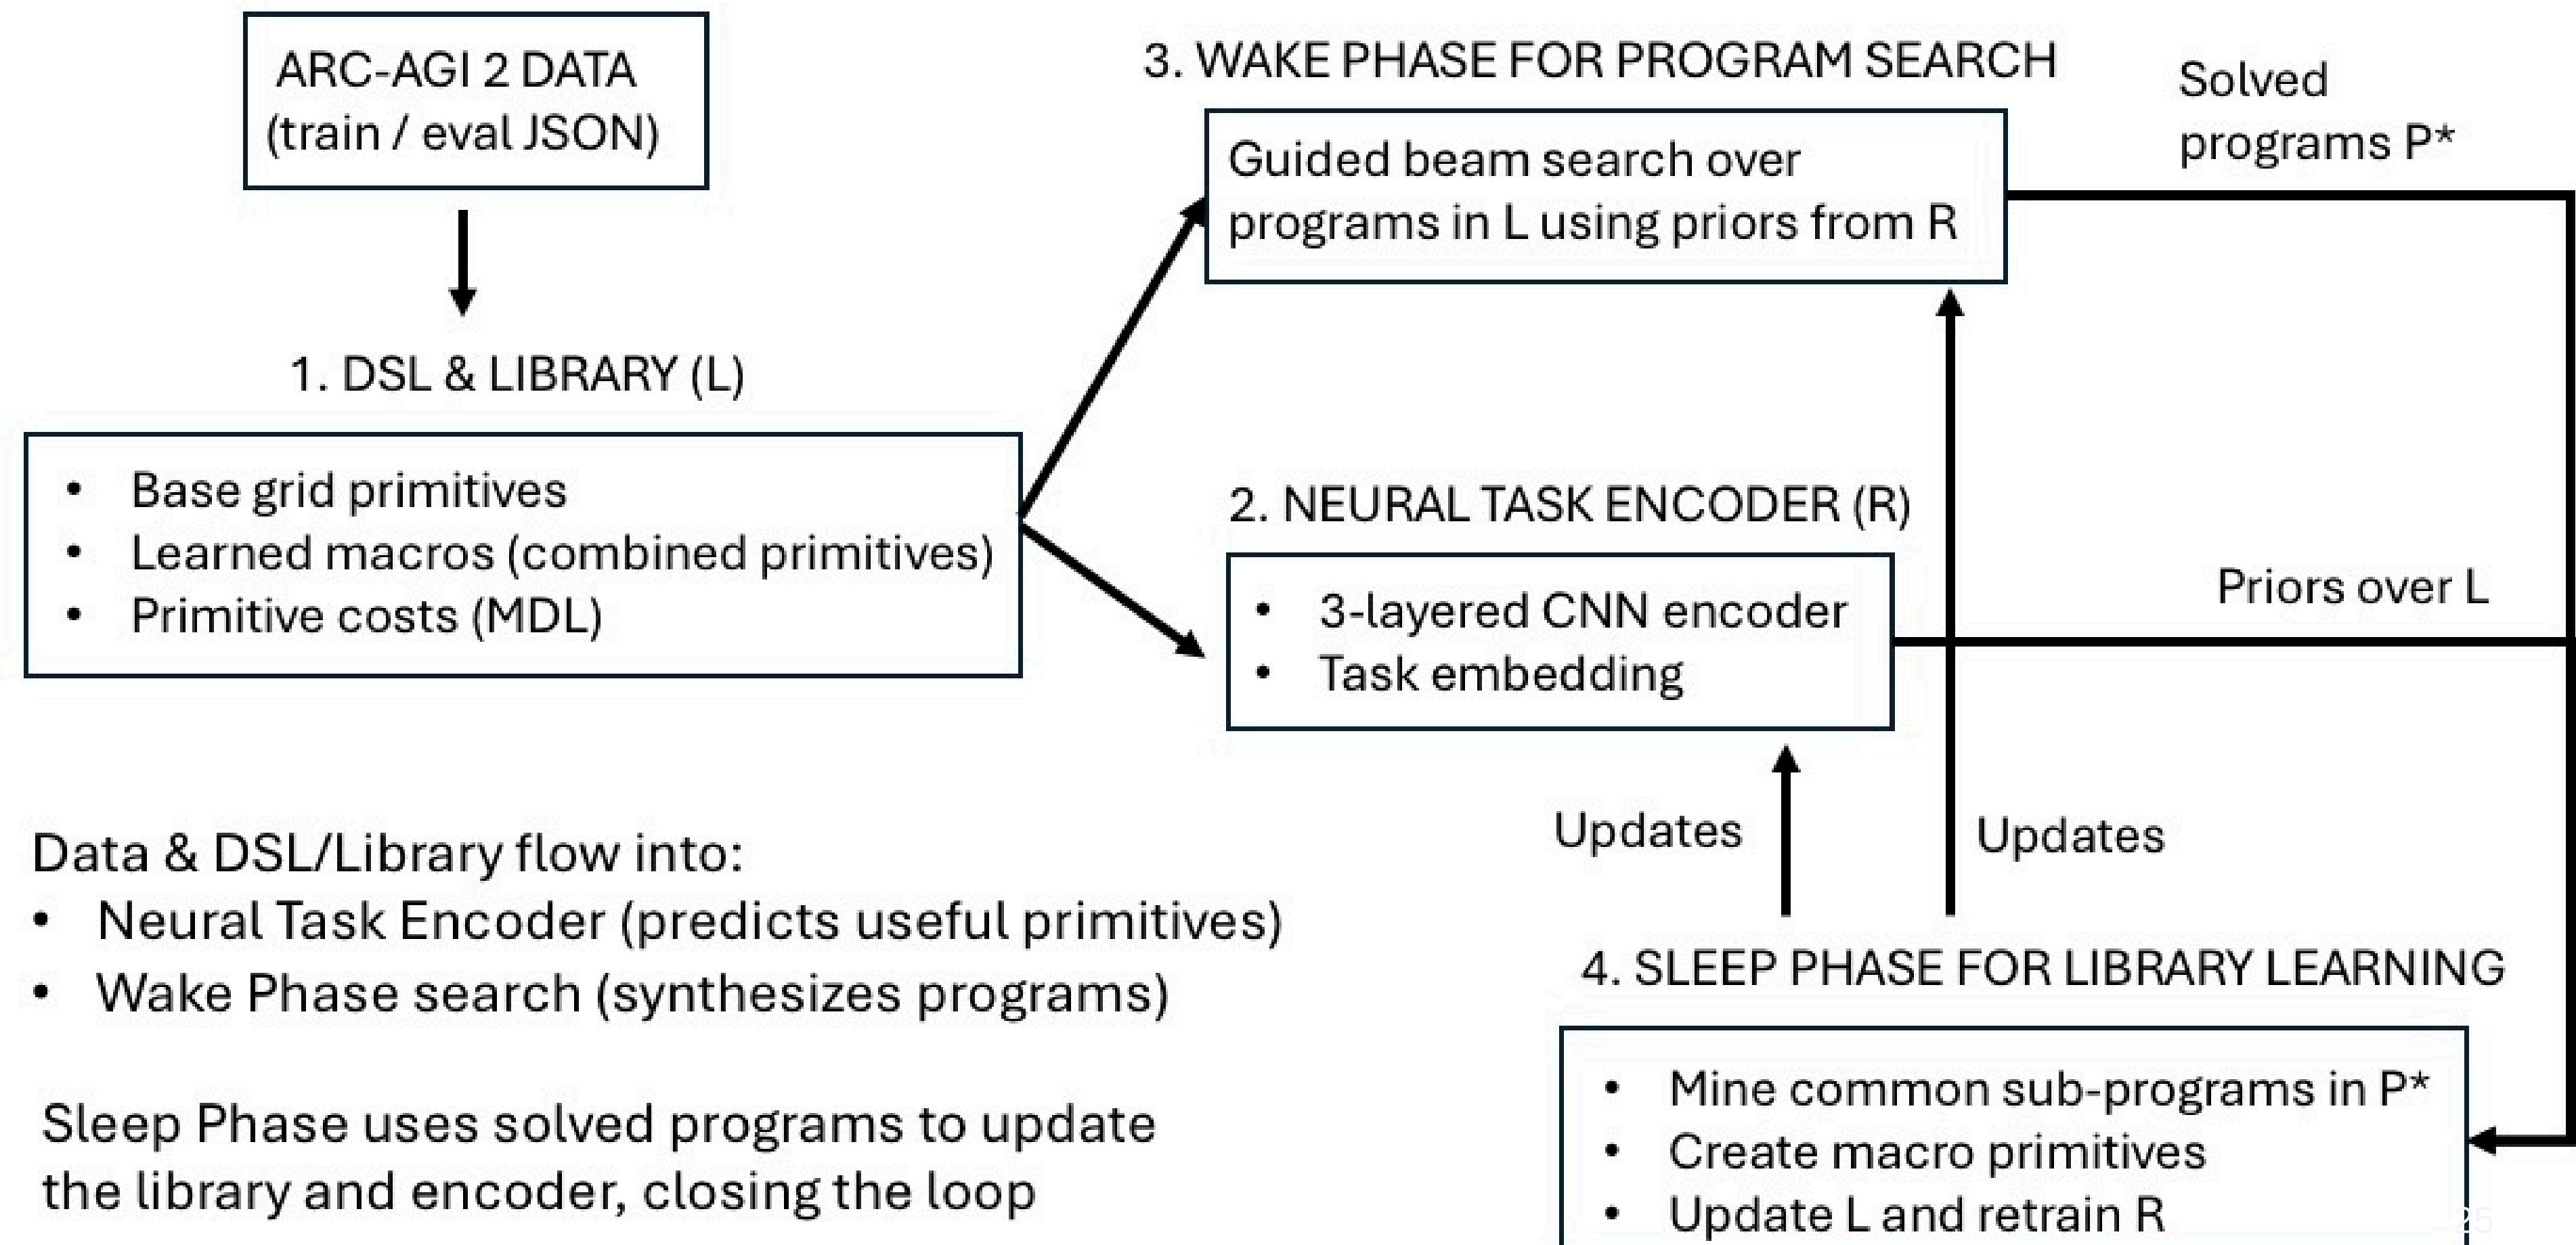
Program-based solutions advantage

- Interpretable (can inspect the synthesized program)
- Reusable (patterns found on one task help on others)

Basic ideas of library learning approach

1. Start with a small DSL
2. Automatically invent higher-level primitives when patterns repeat
3. Judge whether a new primitive is actually useful

Our Architecture Overview



Conclusion

- 1. How to increase number of unique problems solved?**
- 2. How to develop Artificial General Intelligence (AGI)?**

- Create lots of generalized operations (primitives) through code & LLMs
- Retrieving outside context when needed.
- Abstraction to simplify future problem solving.

At every step, evaluate! This informs model which squares are still wrong, helping model learn changes in problems & better adapt

Impact (beyond solving squares)

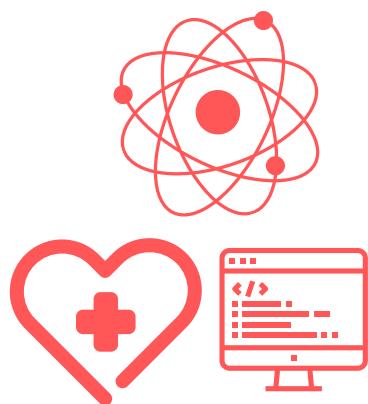
The future for AI is human-reasoning.

Currently, even “PhD-level” AI struggles with tasks outside their trained probability distributions
AI lacks human reasoning skills. Thus, our methodology mimics the best of human reasoning--

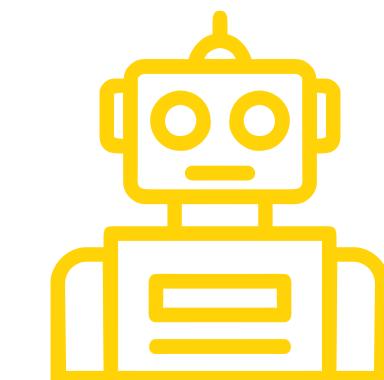
The scientific method:

1. Analyze data & solve empirically - Code primitives
2. Retrieve help from outside - LLM and RAG for context
3. Abstract and generalize - DreamCoder
4. Adapt to changes based on evaluation - Evaluation at every step

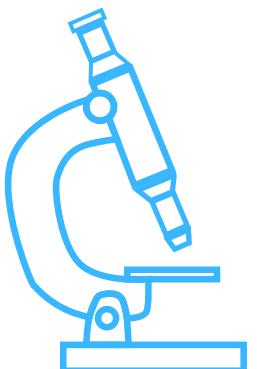
This is how AI can advance to reasoning and general intelligence



Domain-specific applications



AI

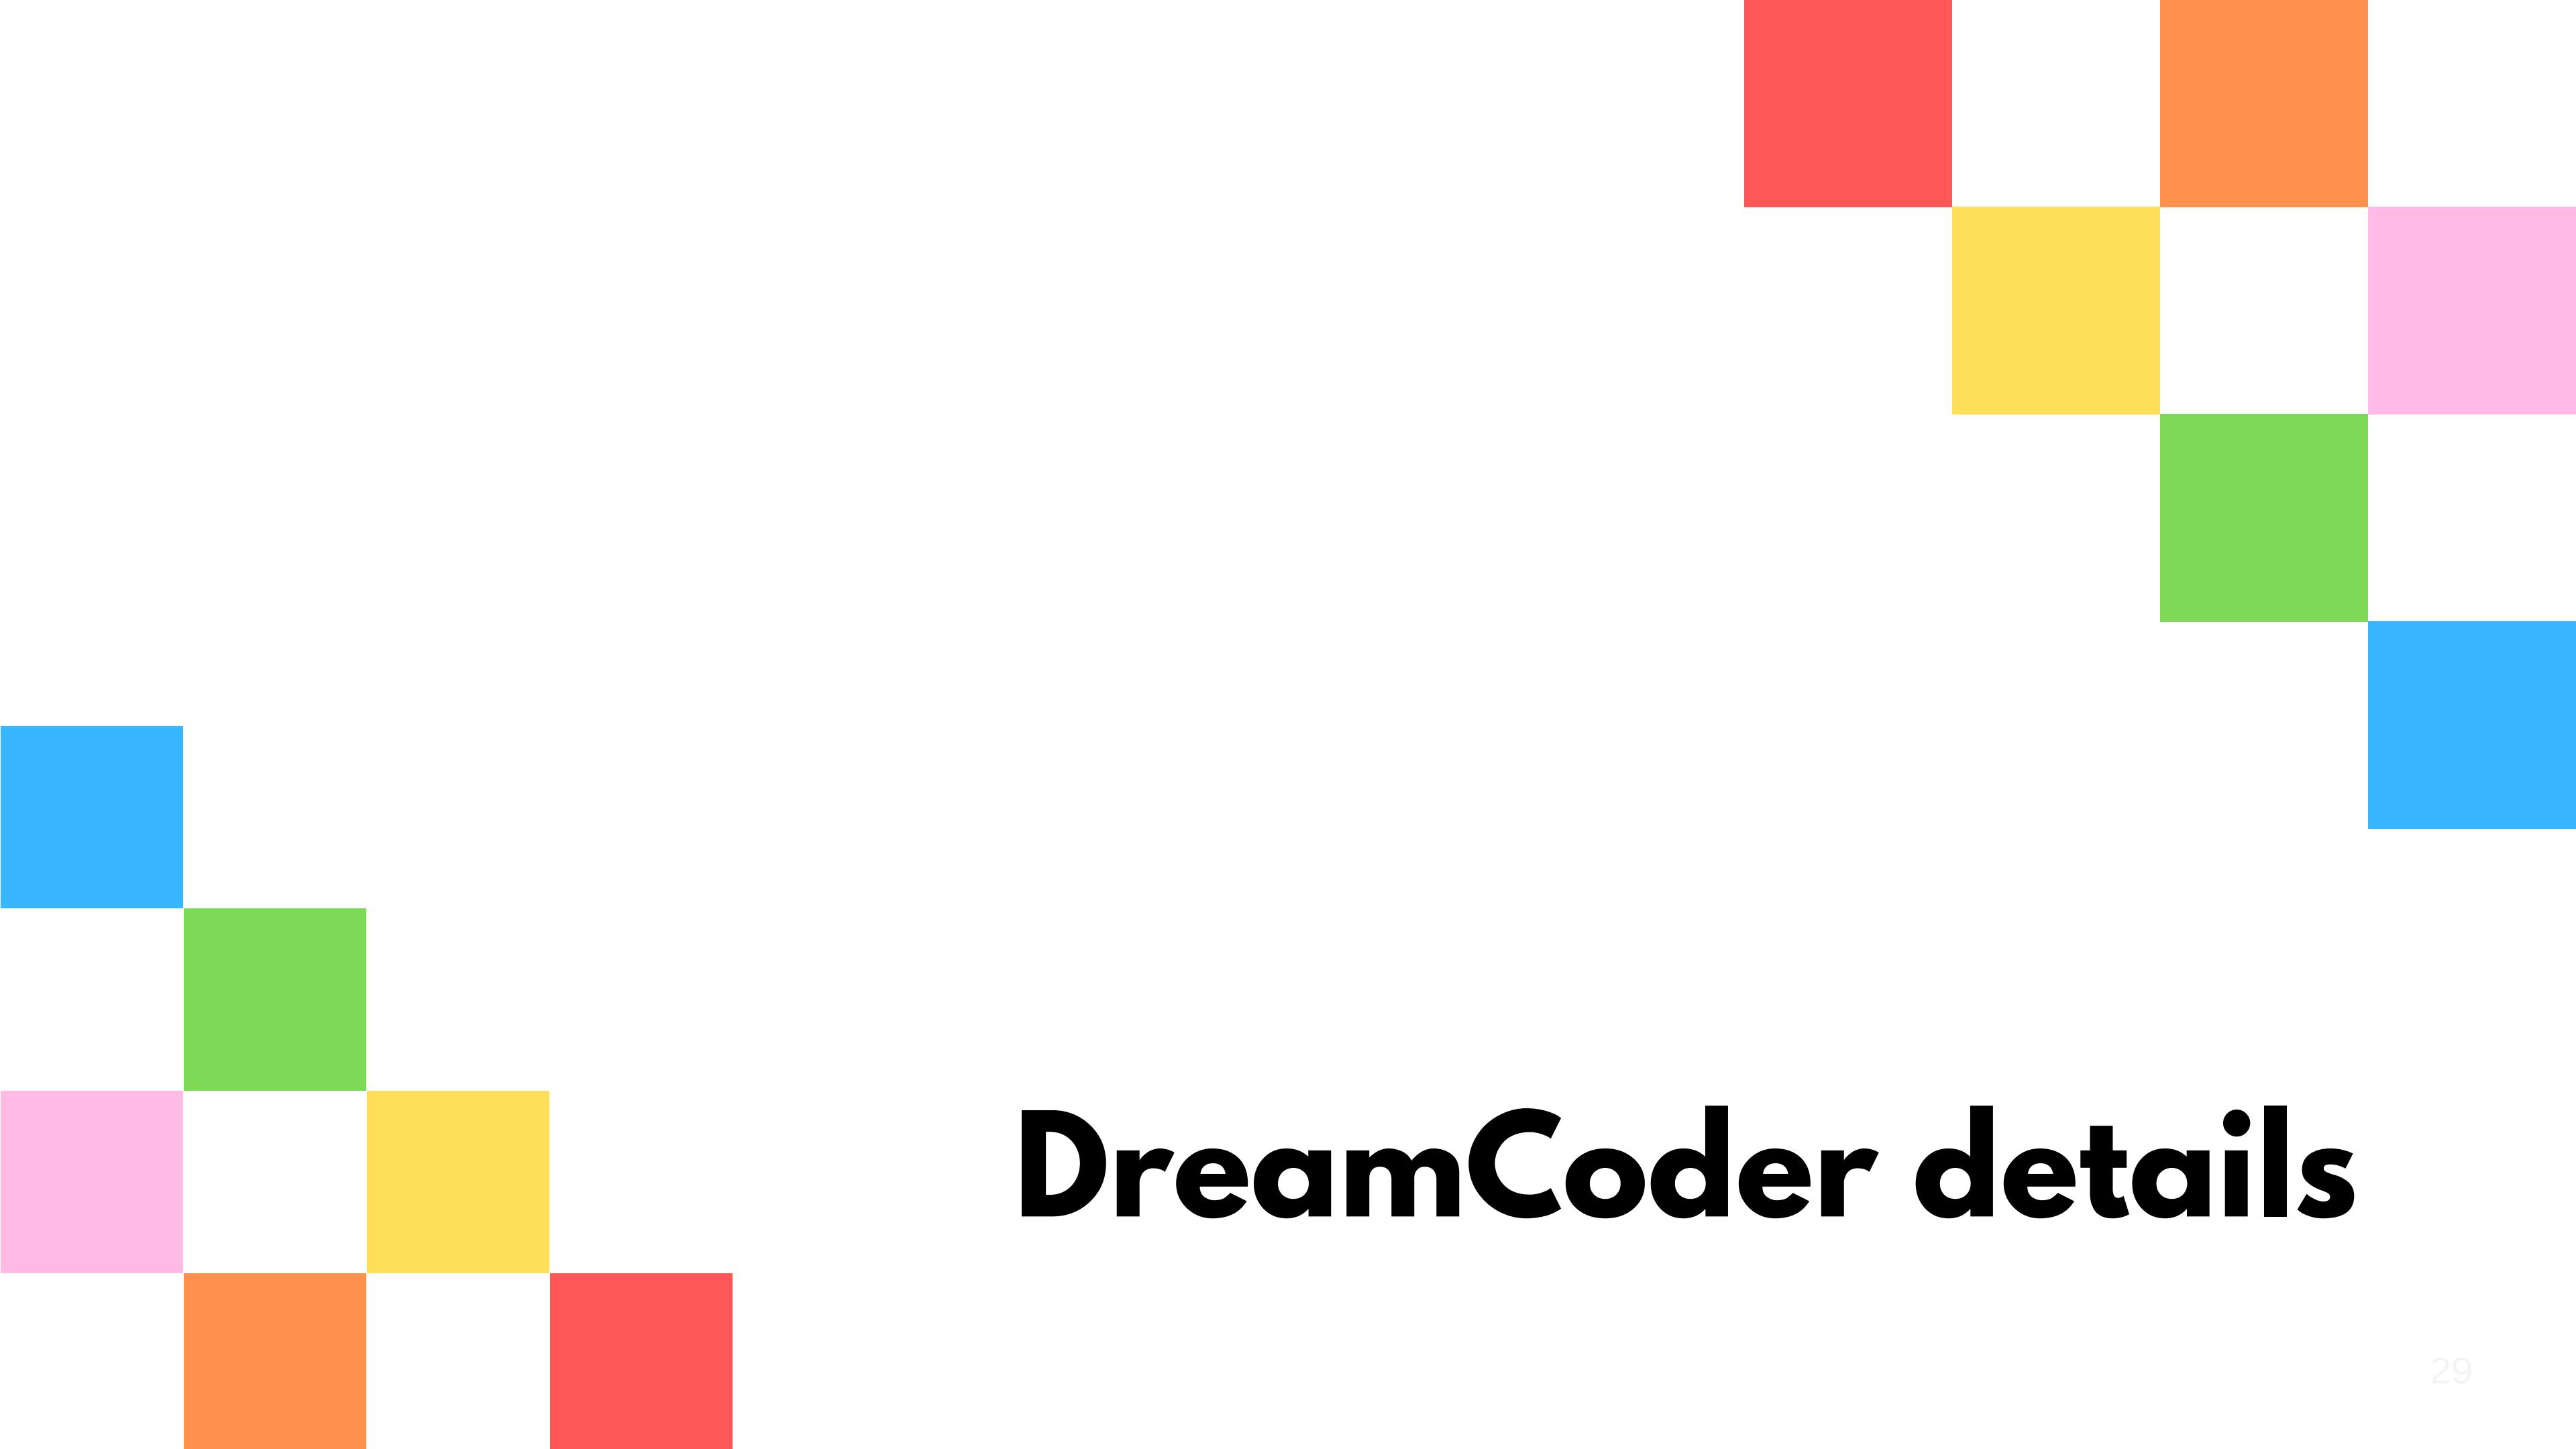


Scientific method



Questions?

Thank you!



DreamCoder details

Architecture part 1: DSL & Library (L)

-The Building Block

Overview

- **Each primitive is implemented as a Python function (mirror, rotate, recolor)**
- **Calculate cost (how many primitives it takes to solve a puzzle)**
- **Reward short & simple programs**

Why?

- **Makes search space smaller**
- **Encourages reuse of general concepts (rotate, crop) instead of task-specific hacks**

Architecture part 2: Neural Task Encoder - Guiding Search

Overview

- **Small CNN that looks at a puzzle and predicts**
- **Curveballs**
 - **Task embedding**
 - **Probability for each primitive in the library**
- **Vector of scores over primitives**

Why?

- **Pure brute-force search over programs is too large**
- **The encoder acts like a heuristic:**
 - **“Try these primitives earlier, those later”**

Architecture part 3: Wake phase - Program Search

Overview

A guided beam search that tries to build a program

Initialize candidate programs

- Start with simple leaf programs using primitives from the library
- Choose parameters (e.g., target sizes, colors) from observed data (grid sizes, color palette)

Why?

- System synthesizes programs that explain the data

Architecture part-3: Sleep phase

Overview

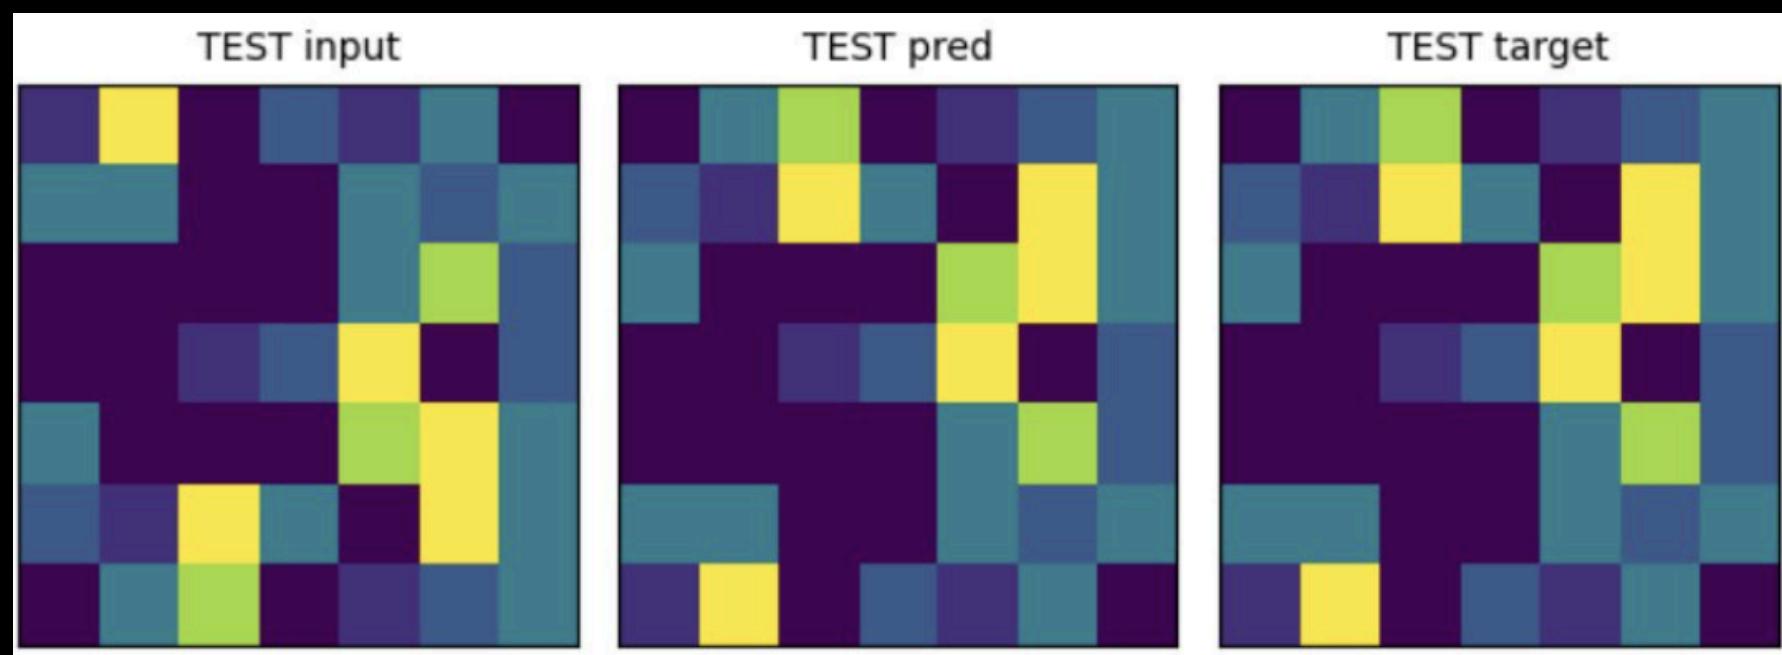
- Collect solved programs and mine frequent sub-programs
- Turn them into new macro primitives with lower Kolmogorov complexity (by program length)
- Update the DSL and retrain the encoder to use these new building blocks

Why?

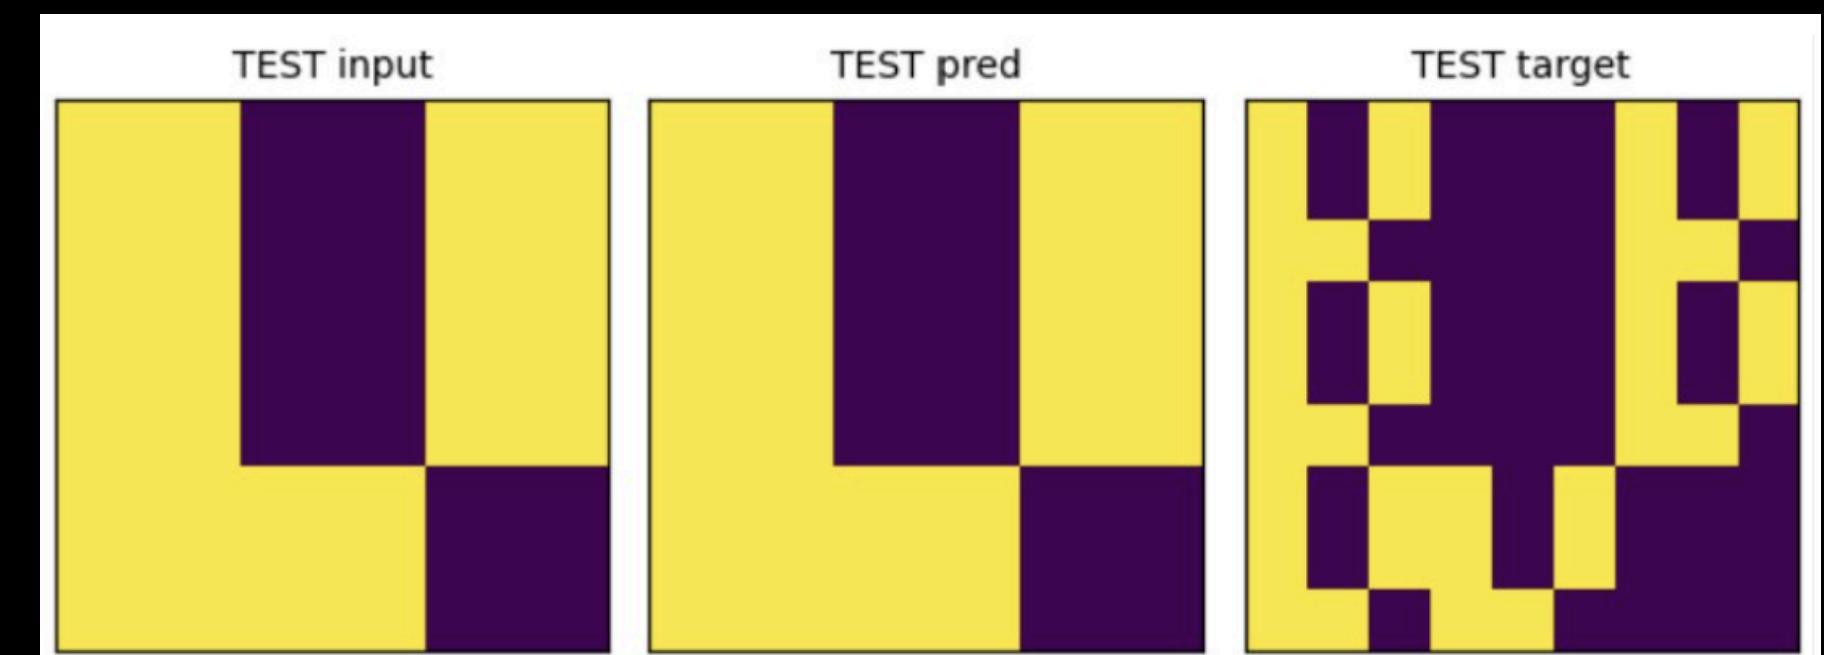
- Over time, the system learns higher-level concepts tailored to the benchmark
- Improves success and learned vocabulary of future search

Results

Correct example

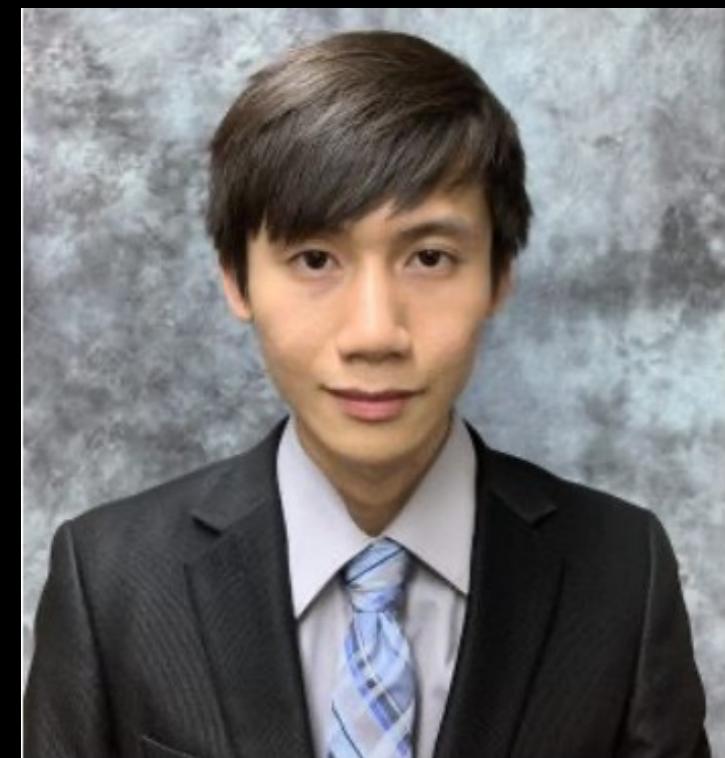


Incorrect example



**Among 300 training samples, the library
managed to solve only 5 puzzles**

Team members



Wen-Hsin Chen | Jeong Wook Lee | Alina Rohulia | Samrat Kumar Dey | Kun-Yi Chen
CS | Informatics | Informatics | Informatics | Informatics